

Personalized cancer diagnosis

1. Business Problem

1.1. Description

Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment/>
(<https://www.kaggle.com/c/msk-redefining-cancer-treatment/>)

Data: Memorial Sloan Kettering Cancer Center (MSKCC)

Download training_variants.zip and training_text.zip from Kaggle.

Context:

Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment/discussion/35336#198462> (<https://www.kaggle.com/c/msk-redefining-cancer-treatment/discussion/35336#198462>)

Problem statement :

Classify the given genetic variations/mutations based on evidence from text-based clinical literature.

1.2. Source/Useful Links

Some articles and reference blogs about the problem statement

1. <https://www.forbes.com/sites/matthewherper/2017/06/03/a-new-cancer-drug-helped-almost-everyone-who-took-it-almost-heres-what-it-teaches-us/#2a44ee2f6b25>
(<https://www.forbes.com/sites/matthewherper/2017/06/03/a-new-cancer-drug-helped-almost-everyone-who-took-it-almost-heres-what-it-teaches-us/#2a44ee2f6b25>)
2. <https://www.youtube.com/watch?v=UwbuW7oK8rk> (<https://www.youtube.com/watch?v=UwbuW7oK8rk>)
3. <https://www.youtube.com/watch?v=qxXRKVompl8> (<https://www.youtube.com/watch?v=qxXRKVompl8>)

1.3. Real-world/Business objectives and constraints.

- No low-latency requirement.
- Interpretability is important.
- Errors can be very costly.
- Probability of a data-point belonging to each class is needed.

2. Machine Learning Problem Formulation

2.1. Data

2.1.1. Data Overview

- Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment/data>
(<https://www.kaggle.com/c/msk-redefining-cancer-treatment/data>)
- We have two data files: one contains the information about the genetic mutations and the other contains the clinical evidence (text) that human experts/pathologists use to classify the genetic mutations.
- Both these data files have a common column called ID
- Data file's information:
 - training_variants (ID , Gene, Variations, Class)
 - training_text (ID, Text)

2.1.2. Example Data Point

training_variants

ID, Gene, Variation, Class
0, FAM58A, Truncating Mutations, 1
1, CBL, W802*, 2
2, CBL, Q249E, 2
...

training_text

ID, Text
0||Cyclin-dependent kinases (CDKs) regulate a variety of fundamental cellular processes. CDK10 stands out as one of the last orphan CDKs for which no activating cyclin has been identified and no kinase activity revealed. Previous work has shown that CDK10 silencing increases ETS2 (v-ets erythroblastosis virus E26 oncogene homolog 2)-driven activation of the MAPK pathway, which confers tamoxifen resistance to breast cancer cells. The precise mechanisms by which CDK10 modulates ETS2 activity, and more generally the functions of CDK10, remain elusive. Here we demonstrate that CDK10 is a cyclin-dependent kinase by identifying cyclin M as an activating cyclin. Cyclin M, an orphan cyclin, is the product of FAM58A, whose mutations cause STAR syndrome, a human developmental anomaly whose features include toe syndactyly, telecanthus, and anogenital and renal malformations. We show that STAR syndrome-associated cyclin M mutants are unable to interact with CDK10. Cyclin M silencing phenocopies CDK10 silencing in increasing c-Raf and in conferring tamoxifen resistance to breast cancer cells. CDK10/cyclin M phosphorylates ETS2 in vitro, and in cells it positively controls ETS2 degradation by the proteasome. ETS2 protein levels are increased in cells derived from a STAR patient, and this increase is attributable to decreased cyclin M levels. Altogether, our results reveal an additional regulatory mechanism for ETS2, which plays key roles in cancer and development. They also shed light on the molecular mechanisms underlying STAR syndrome. Cyclin-dependent kinases (CDKs) play a pivotal role in the control of a number of fundamental cellular processes (1). The human genome contains 21 genes encoding proteins that can be considered as members of the CDK family owing to their sequence similarity with bona fide CDKs, those known to be activated by cyclins (2). Although discovered almost 20 y ago (3, 4), CDK10 remains one of the two CDKs without an identified cyclin partner. This knowledge gap has largely impeded the exploration of its biological functions. CDK10 can act as a positive cell cycle regulator in some cells (5, 6) or as a tumor suppressor in others (7, 8). CDK10 interacts with the ETS2 (v-ets erythroblastosis virus E26 oncogene homolog 2) transcription factor and inhibits its transcriptional activity through an unknown mechanism (9). CDK10 knockdown derepresses ETS2, which increases the expression of the c-Raf protein kinase, activates the MAPK pathway, and induces resistance of MCF7 cells to tamoxifen (6). ...

2.2. Mapping the real-world problem to an ML problem

2.2.1. Type of Machine Learning Problem

There are nine different classes a genetic mutation can be classified into => Multi class classification problem

2.2.2. Performance Metric

Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment#evaluation>
(<https://www.kaggle.com/c/msk-redefining-cancer-treatment#evaluation>).

Metric(s):

- Multi class log-loss
- Confusion matrix

2.2.3. Machine Learning Objectives and Constraints

Objective: Predict the probability of each data-point belonging to each of the nine classes.

Constraints:

- Interpretability
- Class probabilities are needed.
- Penalize the errors in class probabilities => Metric is Log-loss.
- No Latency constraints.

2.3. Train, CV and Test Datasets

Split the dataset randomly into three parts train, cross validation and test with 64%, 16%, 20% of data respectively

3. Exploratory Data Analysis

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import SGDClassifier
from imblearn.over_sampling import SMOTE
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_extraction.text import TfidfVectorizer
warnings.filterwarnings("ignore")

from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
```

3.1. Reading Data

3.1.1. Reading Gene and Variation Data

```
In [ ]: data = pd.read_csv('training_variants')
print('Number of data points : ', data.shape[0])
print('Number of features : ', data.shape[1])
print('Features : ', data.columns.values)
data.head()
```

training/training_variants is a comma separated file containing the description of the genetic mutations used for training.

Fields are

- **ID** : the id of the row used to link the mutation to the clinical evidence
- **Gene** : the gene where this genetic mutation is located
- **Variation** : the aminoacid change for this mutations
- **Class** : 1-9 the class this genetic mutation has been classified on

3.1.2. Reading Text Data

```
In [3]: # note the separator in this file
data_text = pd.read_csv("training_text", sep="\\|\\|", engine="python", names=
print('Number of data points : ', data_text.shape[0])
print('Number of features : ', data_text.shape[1])
print('Features : ', data_text.columns.values)
data_text.head()
```

```
Number of data points : 3321
Number of features : 2
Features : ['ID' 'TEXT']
```

Out[3]:

	ID	TEXT
0	0	Cyclin-dependent kinases (CDKs) regulate a var...
1	1	Abstract Background Non-small cell lung canc...
2	2	Abstract Background Non-small cell lung canc...
3	3	Recent evidence has demonstrated that acquired...
4	4	Oncogenic mutations in the monomeric Casitas B...

3.1.3. Preprocessing of text

```
In [4]: # loading stop words from nltk library
stop_words = set(stopwords.words('english'))

def nlp_preprocessing(total_text, index, column):
    if type(total_text) is not int:
        string = ""
        # replace every special char with space
        total_text = re.sub('[^a-zA-Z0-9\n]', ' ', total_text)
        # replace multiple spaces with single space
        total_text = re.sub('\s+', ' ', total_text)
        # converting all the chars into lower-case.
        total_text = total_text.lower()

        for word in total_text.split():
            # if the word is a not a stop word then retain that word from
            if not word in stop_words:
                string += word + " "

        data_text[column][index] = string
```

```
In [5]: #text processing stage.
start_time = time.clock()
for index, row in data_text.iterrows():
    if type(row['TEXT']) is str:
        nlp_preprocessing(row['TEXT'], index, 'TEXT')
    else:
        print("there is no text description for id:",index)
print('Time took for preprocessing the text :',time.clock() - start_time)
```

```
there is no text description for id: 1109
there is no text description for id: 1277
there is no text description for id: 1407
there is no text description for id: 1639
there is no text description for id: 2755
Time took for preprocessing the text : 225.53517599999998 seconds
```

```
In [6]: #merging both gene_variations and text data based on ID
result = pd.merge(data, data_text,on='ID', how='left')
result.head()
```

Out[6]:

	ID	Gene	Variation	Class	TEXT
0	0	FAM58A	Truncating Mutations	1	cyclin dependent kinases cdks regulate variety...
1	1	CBL	W802*	2	abstract background non small cell lung cancer...
2	2	CBL	Q249E	2	abstract background non small cell lung cancer...
3	3	CBL	N454D	3	recent evidence demonstrated acquired uniparen...
4	4	CBL	L399V	4	oncogenic mutations monomeric casitas b lineag...

```
In [7]: result[result.isnull().any(axis=1)]
```

Out[7]:

	ID	Gene	Variation	Class	TEXT
1109	1109	FANCA	S1088F	1	NaN
1277	1277	ARID5B	Truncating Mutations	1	NaN
1407	1407	FGFR3	K508M	6	NaN
1639	1639	FLT1	Amplification	6	NaN
2755	2755	BRAF	G596C	7	NaN

```
In [8]: result.loc[result['TEXT'].isnull(), 'TEXT'] = result['Gene'] + ' ' + result['Variation']
```

```
In [9]: result[result['ID']==1109]
```

Out[9]:

	ID	Gene	Variation	Class	TEXT
1109	1109	FANCA	S1088F	1	FANCA S1088F

Feature Engineering of given data

```
In [10]: # number of words in each column
result["no_of_words"] = result["TEXT"].apply(lambda x: len(x.split()))
result.head()
```

Out[10]:

	ID	Gene	Variation	Class	TEXT	no_of_words
0	0	FAM58A	Truncating Mutations	1	cyclin dependent kinases cdks regulate variety...	4370
1	1	CBL	W802*	2	abstract background non small cell lung cancer...	4139
2	2	CBL	Q249E	2	abstract background non small cell lung cancer...	4139
3	3	CBL	N454D	3	recent evidence demonstrated acquired uniparen...	3841
4	4	CBL	L399V	4	oncogenic mutations monomeric casitas b lineag...	4254


```
In [11]: # number of characters in the column
result['no_of_characters'] = result['TEXT'].apply(lambda x: len(str(x)))
result.head()
```

Out[11]:

	ID	Gene	Variation	Class	TEXT	no_of_words	no_of_characters
0	0	FAM58A	Truncating Mutations	1	cyclin dependent kinases cdks regulate variety...	4370	30836
1	1	CBL	W802*	2	abstract background non small cell lung cancer...	4139	27844
2	2	CBL	Q249E	2	abstract background non small cell lung cancer...	4139	27844
3	3	CBL	N454D	3	recent evidence demonstrated acquired uniparen...	3841	28093
4	4	CBL	L399V	4	oncogenic mutations monomeric casitas b lineag...	4254	31649


```
In [12]: # combining the Gene column and Variation column
result['gene_variation'] = result['Gene'] + " " + result["Variation"]
result.head()
```

Out[12]:

	ID	Gene	Variation	Class	TEXT	no_of_words	no_of_characters	gene_variation
0	0	FAM58A	Truncating Mutations	1	cyclin dependent kinases cdks regulate variety...	4370	30836	FAM58A Truncating Mutation
1	1	CBL	W802*	2	abstract background non small cell lung cancer...	4139	27844	CBL W802*
2	2	CBL	Q249E	2	abstract background non small cell lung cancer...	4139	27844	CBL Q249E
3	3	CBL	N454D	3	recent evidence demonstrated acquired uniparen...	3841	28093	CBL N454D
4	4	CBL	L399V	4	oncogenic mutations monomeric casitas b lineag...	4254	31649	CBL L399V

```
In [13]: plt.hist(result["no_of_words"], normed=True, bins=100)
plt.xlabel('Counts');
plt.title("Count of number of words")
```

```
Out[13]: Text(0.5,1,'Count of number of words')
```



```
In [14]: # we can also build a column of if the no. of words are greater than 5
result["word_count_of_5k"] = result["no_of_words"].apply(lambda x: 1 if x > 5 else 0)
print("Head result",result.head(5))
print("*****")
print("\nTail result",result.tail())
```

	Head result	ID	Gene	Variation	Class \
0	0	FAM58A	Truncating Mutations	1	
1	1	CBL	W802*	2	
2	2	CBL	Q249E	2	
3	3	CBL	N454D	3	
4	4	CBL	L399V	4	

	TEXT	no_of_words \
0	cyclin dependent kinases cdks regulate variety...	4370
1	abstract background non small cell lung cancer...	4139
2	abstract background non small cell lung cancer...	4139
3	recent evidence demonstrated acquired uniparen...	3841
4	oncogenic mutations monomeric casitas b lineag...	4254

	no_of_characters	gene_variation	word_count_of_5k
0	30836	FAM58A Truncating Mutations	0
1	27844	CBL W802*	0
2	27844	CBL Q249E	0
3	28093	CBL N454D	0
4	31649	CBL L399V	0


	Tail result	ID	Gene	Variation	Class \
3316	3316	RUNX1	D171N	4	
3317	3317	RUNX1	A122*	1	
3318	3318	RUNX1	Fusions	1	
3319	3319	RUNX1	R80C	4	
3320	3320	RUNX1	K83E	4	

	TEXT	no_of_words
3316	introduction myelodysplastic syndromes mds het...	8153
3317	introduction myelodysplastic syndromes mds het...	4495
3318	runt related transcription factor 1 gene runx1...	4593
3319	runx1 aml1 gene frequent target chromosomal tr...	3465
3320	frequent mutations associated leukemia recurrence...	7013

	no_of_characters	gene_variation	word_count_of_5k
3316	57217	RUNX1 D171N	1
3317	30800	RUNX1 A122*	0
3318	28080	RUNX1 Fusions	0
3319	25404	RUNX1 R80C	0
3320	52582	RUNX1 K83E	1

```
In [15]: plt.hist(result["no_of_characters"], normed=True, bins=100)
plt.xlabel('Counts');
plt.title("Count of number of characters")
```

```
Out[15]: Text(0.5,1,'Count of number of characters')
```



```
In [16]: # we can also build a column of if the no. of characters are greater than 50k
result["character_count_of_50k"] = result["no_of_characters"].apply(lambda x: 1 if x > 50000 else 0)
print("Head result",result.head(5))
print("*****")
print("\nTail result",result.tail())
```

	Head result	ID	Gene	Variation	Class \
0	0	FAM58A	Truncating Mutations	1	
1	1	CBL	W802*	2	
2	2	CBL	Q249E	2	
3	3	CBL	N454D	3	
4	4	CBL	L399V	4	

	TEXT	no_of_words \
0	cyclin dependent kinases cdks regulate variety...	4370
1	abstract background non small cell lung cancer...	4139
2	abstract background non small cell lung cancer...	4139
3	recent evidence demonstrated acquired uniparen...	3841
4	oncogenic mutations monomeric casitas b lineag...	4254

	no_of_characters	gene_variation	word_count_of_5k
0	30836	FAM58A Truncating Mutations	0
1	27844	CBL W802*	0
2	27844	CBL Q249E	0
3	28093	CBL N454D	0
4	31649	CBL L399V	0

	character_count_of_50k
0	0

```

1          0
2          0
3          0
4          0
*****

```

```

Tail result      ID   Gene Variation  Class \
3316  3316  RUNX1      D171N      4
3317  3317  RUNX1      A122*      1
3318  3318  RUNX1      Fusions     1
3319  3319  RUNX1      R80C       4
3320  3320  RUNX1      K83E       4

```

```

\
TEXT  no_of_words
3316  introduction myelodysplastic syndromes mds het... 8153
3317  introduction myelodysplastic syndromes mds het... 4495
3318  runt related transcription factor 1 gene runx1... 4593
3319  runx1 aml1 gene frequent target chromosomal tr... 3465
3320  frequent mutations associated leukemia recurrence... 7013

```

```

no_of_characters gene_variation  word_count_of_5k \
3316          57217      RUNX1 D171N          1
3317          30800      RUNX1 A122*          0
3318          28080      RUNX1 Fusions          0
3319          25404      RUNX1 R80C           0
3320          52582      RUNX1 K83E           1

```

```

character_count_of_50k
3316          1
3317          0
3318          0
3319          0
3320          1

```

```
In [17]: result.columns.values
```

```
Out[17]: array(['ID', 'Gene', 'Variation', 'Class', 'TEXT', 'no_of_words',
                'no_of_characters', 'gene_variation', 'word_count_of_5k',
                'character_count_of_50k'], dtype=object)
```

3.1.4. Test, Train and Cross Validation Split

3.1.4.1. Splitting data into train, test and cross validation (64:20:16)

```
In [18]: y_true = result['Class'].values
result.Gene      = result.Gene.str.replace('\s+', '_')
result.Variation = result.Variation.str.replace('\s+', '_')

# split the data into test and train by maintaining same distribution
X_train, test_df, y_train, y_test = train_test_split(result, y_true, s
# split the train data into train and cross validation by maintaining
train_df, cv_df, y_train, y_cv = train_test_split(X_train, y_train, st
```

We split the data into train, test and cross validation data sets, preserving the ratio of class distribution in the original data set

```
In [19]: print('Number of data points in train data:', train_df.shape[0])
print('Number of data points in test data:', test_df.shape[0])
print('Number of data points in cross validation data:', cv_df.shape[0])

Number of data points in train data: 2124
Number of data points in test data: 665
Number of data points in cross validation data: 532
```

3.1.4.2. Distribution of y_i's in Train, Test and Cross Validation datasets

```
In [20]: # it returns a dict, keys as class labels and values as the number of
train_class_distribution = train_df['Class'].value_counts().sortlevel()
test_class_distribution = test_df['Class'].value_counts().sortlevel()
cv_class_distribution = cv_df['Class'].value_counts().sortlevel()

my_colors = 'rgbkymc'
train_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in train data')
plt.grid()
plt.show()


# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
# -(train_class_distribution.values): the minus sign will give us in d
sorted_yi = np.argsort(-train_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ':', train_class_distr

print('-'*80)
my_colors = 'rgbkymc'
test_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in test data')
plt.grid()
plt.show()
```


```
# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
# -(train_class_distribution.values): the minus sign will give us in descending order
sorted_yi = np.argsort(-test_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ': ', test_class_distribution.values[i])


print('-'*80)
my_colors = 'rgbkymc'
cv_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in cross validation data')
plt.grid()
plt.show()

# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
# -(train_class_distribution.values): the minus sign will give us in descending order
sorted_yi = np.argsort(-train_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ': ', cv_class_distribution.values[i])
```




```
Number of data points in class 7 : 609 ( 28.672 %)
Number of data points in class 4 : 439 ( 20.669 %)
Number of data points in class 1 : 363 ( 17.09 %)
Number of data points in class 2 : 289 ( 13.606 %)
Number of data points in class 6 : 176 ( 8.286 %)
Number of data points in class 5 : 155 ( 7.298 %)
Number of data points in class 3 : 57 ( 2.684 %)
Number of data points in class 9 : 24 ( 1.13 %)
Number of data points in class 8 : 12 ( 0.565 %)
```





Number of data points in class 7 : 191 (28.722 %)
 Number of data points in class 4 : 137 (20.602 %)
 Number of data points in class 1 : 114 (17.143 %)
 Number of data points in class 2 : 91 (13.684 %)
 Number of data points in class 6 : 55 (8.271 %)
 Number of data points in class 5 : 48 (7.218 %)
 Number of data points in class 3 : 18 (2.707 %)
 Number of data points in class 9 : 7 (1.053 %)
 Number of data points in class 8 : 4 (0.602 %)



Number of data points in class 7 : 153 (28.759 %)
 Number of data points in class 4 : 110 (20.677 %)
 Number of data points in class 1 : 91 (17.105 %)
 Number of data points in class 2 : 72 (13.534 %)
 Number of data points in class 6 : 44 (8.271 %)
 Number of data points in class 5 : 39 (7.331 %)
 Number of data points in class 3 : 14 (2.632 %)
 Number of data points in class 9 : 6 (1.128 %)
 Number of data points in class 8 : 3 (0.564 %)

3.2 Prediction using a 'Random' Model

In a 'Random' Model, we generate the NINE class probabilities randomly such that they sum to 1.

```
In [21]: # This function plots the confusion matrices given y_i, y_i_hat.
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    # C = 9,9 matrix, each cell (i,j) represents number of points of class i predicted as class j

    A = ((C.T)/(C.sum(axis=1))).T
    #divid each element of the confusion matrix with the sum of elements in each column

    # C = [[1, 2],
    #      [3, 4]]
    # C.T = [[1, 3],
    #        [2, 4]]
    # C.sum(axis = 1) axis=0 corresponds to columns and axis=1 corresponds to rows
    # C.sum(axis = 1) = [[3, 7]]
    # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7],
    #                             [2/3, 4/7]]

    # ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3],
    #                               [3/7, 4/7]]
    # sum of row elements = 1

    B = (C/C.sum(axis=0))
    #divid each element of the confusion matrix with the sum of elements in each row

    # C = [[1, 2],
    #      [3, 4]]
    # C.sum(axis = 0) axis=0 corresponds to columns and axis=1 corresponds to rows
    # C.sum(axis = 0) = [[4, 6]]
    # (C/C.sum(axis=0)) = [[1/4, 2/6],
    #                       [3/4, 4/6]]

    labels = [1,2,3,4,5,6,7,8,9]
    # representing A in heatmap format
    print("-"*20, "Confusion matrix", "-"*20)
    plt.figure(figsize=(20,7))
    sns.heatmap(C, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.show()

    print("-"*20, "Precision matrix (Column Sum=1)", "-"*20)
    plt.figure(figsize=(20,7))
    sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.show()
```

```
# representing B in heatmap format
print("-"*20, "Recall matrix (Row sum=1)", "-"*20)
plt.figure(figsize=(20,7))
sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=1)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```

```
In [22]: # we need to generate 9 numbers and the sum of numbers should be 1
# one solution is to generate 9 numbers and divide each of the numbers
# ref: https://stackoverflow.com/a/18662466/4084039
test_data_len = test_df.shape[0]
cv_data_len = cv_df.shape[0]

# we create a output array that has exactly same size as the CV data
cv_predicted_y = np.zeros((cv_data_len,9))
for i in range(cv_data_len):
    rand_probs = np.random.rand(1,9)
    cv_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Cross Validation Data using Random Model",log_loss(y_cv,cv_predicted_y))


# Test-Set error.
#we create a output array that has exactly same as the test data
test_predicted_y = np.zeros((test_data_len,9))
for i in range(test_data_len):
    rand_probs = np.random.rand(1,9)
    test_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Test Data using Random Model",log_loss(y_test,test_predicted_y))

predicted_y =np.argmax(test_predicted_y, axis=1)
plot_confusion_matrix(y_test, predicted_y+1)
```



Log loss on Cross Validation Data using Random Model 2.4640965314393575

Log loss on Test Data using Random Model 2.446975397398525

----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



3.3 Univariate Analysis

```
In [23]: # code for response coding with Laplace smoothing.
# alpha : used for laplace smoothing
# feature: ['gene', 'variation']
# df: ['train_df', 'test_df', 'cv_df']
# algorithm
# -----
# Consider all unique values and the number of occurrences of given fea
# build a vector (1*9) , the first element = (number of times it occur
# gv_dict is like a look up table, for every gene it store a (1*9) rep
# for a value of feature in df:
# if it is in train data:
# we add the vector that was stored in 'gv_dict' look up table to 'gv_
# if it is not there is train:
# we add [1/9, 1/9, 1/9, 1/9, 1/9, 1/9, 1/9, 1/9] to 'gv_fea'
# return 'gv_fea'
# -----
```

```

# get_gv_fea_dict: Get Gene variation Feature Dict
def get_gv_fea_dict(alpha, feature, df):
    # value_count: it contains a dict like
    # print(train_df['Gene'].value_counts())
    # output:
    #      {BRCA1      174
    #       TP53      106
    #       EGFR       86
    #       BRCA2       75
    #       PTEN       69
    #       KIT        61
    #       BRAF        60
    #       ERBB2       47
    #       PDGFRA      46
    #       ...}
    # print(train_df['Variation'].value_counts())
    # output:
    # {
    # Truncating_Mutations      63
    # Deletion                  43
    # Amplification              43
    # Fusions                   22
    # Overexpression             3
    # E17K                      3
    # Q61L                      3
    # S222D                      2
    # P130S                      2
    # ...
    # }
    value_count = train_df[feature].value_counts()

    # gv_dict : Gene Variation Dict, which contains the probability array
    gv_dict = dict()

    # denominator will contain the number of time that particular feature occurs
    for i, denominator in value_count.items():
        # vec will contain (p(yi=1/Gi) probability of gene/variation occurrence)
        # vec is 9 dimensional vector
        vec = []
        for k in range(1,10):
            # print(train_df.loc[(train_df['Class']==1) & (train_df['Gene']==i)])
            #      ID      Gene      Variation      Class
            # 2470  2470  BRCA1      S1715C      1
            # 2486  2486  BRCA1      S1841R      1
            # 2614  2614  BRCA1      M1R        1
            # 2432  2432  BRCA1      L1657P      1
            # 2567  2567  BRCA1      T1685A      1
            # 2583  2583  BRCA1      E1660G      1
            # 2634  2634  BRCA1      W1718L      1
            # cls_cnt.shape[0] will return the number of rows

            cls_cnt = train_df.loc[(train_df['Class']==k) & (train_df['Gene']==i)]

            # cls_cnt.shape[0](numerator) will contain the number of times

```

```

        vec.append((cls_cnt.shape[0] + alpha*10) / (denominator + 90))

    # we are adding the gene/variation to the dict as key and vec as value
    gv_dict[i]=vec
return gv_dict

# Get Gene variation feature
def get_gv_feature(alpha, feature, df):
    # print(gv_dict)
    # { 'BRCA1': [0.20075757575757575, 0.03787878787878788, 0.06818181818181818],
    #    'TP53': [0.32142857142857145, 0.061224489795918366, 0.061224489795918366],
    #    'EGFR': [0.056818181818181816, 0.21590909090909091, 0.0625, 0.0625, 0.0625],
    #    'BRCA2': [0.13333333333333333, 0.060606060606060608, 0.060606060606060608, 0.060606060606060608],
    #    'PTEN': [0.069182389937106917, 0.062893081761006289, 0.069182389937106917, 0.069182389937106917],
    #    'KIT': [0.066225165562913912, 0.25165562913907286, 0.07284768211917995, 0.07284768211917995],
    #    'BRAF': [0.066666666666666666, 0.17999999999999999, 0.07333333333333333, 0.07333333333333333],
    #    ...
    # }
    gv_dict = get_gv_fea_dict(alpha, feature, df)
    # value_count is similar in get_gv_fea_dict
    value_count = train_df[feature].value_counts()

    # gv_fea: Gene_variation feature, it will contain the feature for each row
    gv_fea = []
    # for every feature values in the given data frame we will check if it is in the dict
    # if not we will add [1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9] to gv_fea
    for index, row in df.iterrows():
        if row[feature] in dict(value_count).keys():
            gv_fea.append(gv_dict[row[feature]])
        else:
            gv_fea.append([1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9])
    # gv_fea.append([-1,-1,-1,-1,-1,-1,-1,-1,-1])
return gv_fea

```

when we calculate the probability of a feature belongs to any particular class, we apply laplace smoothing

- $(\text{numerator} + 10 \cdot \alpha) / (\text{denominator} + 90 \cdot \alpha)$

3.2.1 Univariate Analysis on Gene Feature

Q1. Gene, What type of feature it is ?

Ans. Gene is a categorical variable

Q2. How many categories are there and How they are distributed?


```
In [24]: unique_genes = train_df['Gene'].value_counts()
print('Number of Unique Genes :', unique_genes.shape[0])
# the top 10 genes that occurred most
print(unique_genes.head(10))
```

```
Number of Unique Genes : 224
BRCA1      176
TP53       100
EGFR        95
BRCA2       85
PTEN        81
KIT          63
BRAF        55
ERBB2       39
ALK          38
FGFR2       34
Name: Gene, dtype: int64
```


```
In [25]: print("Ans: There are", unique_genes.shape[0] , "different categories o
```

Ans: There are 224 different categories of genes in the train data, and they are distributed as follows

```
In [26]: s = sum(unique_genes.values);
h = unique_genes.values/s;
plt.plot(h, label="Histogram of Genes")
plt.xlabel('Index of a Gene')
plt.ylabel('Number of Occurances')
plt.legend()
plt.grid()
plt.show()
```



```
In [27]: c = np.cumsum(h)
plt.plot(c, label='Cumulative distribution of Genes')
plt.grid()
plt.legend()
plt.show()
```



Q3. How to featurize this Gene feature ?

Ans. there are two ways we can featurize this variable check out this video:

<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>
[\(https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/\)](https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/)

1. One hot Encoding
2. Response coding

We will choose the appropriate featurization based on the ML model we use. For this problem of multi-class classification with categorical features, one-hot encoding is better for Logistic regression while response coding is better for Random Forests.

```
In [28]: #response-coding of the Gene feature
# alpha is used for laplace smoothing
alpha = 1
# train gene feature
train_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene"))
# test gene feature
test_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene"))
# cross validation gene feature
cv_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene"))
```

```
In [29]: print("train_gene_feature_responseCoding is converted feature using re  

train_gene_feature_responseCoding is converted feature using response  

coding method. The shape of gene feature: (2124, 9)
```

```
In [30]: # one-hot encoding of Gene feature.  

gene_vectorizer = CountVectorizer(max_features=220)  

train_gene_feature_onehotCoding = gene_vectorizer.fit_transform(train_  

test_gene_feature_onehotCoding = gene_vectorizer.transform(test_df['Ge  

cv_gene_feature_onehotCoding = gene_vectorizer.transform(cv_df['Gene']
```

```
In [31]: train_gene_feature_onehotCoding.shape
```

```
Out[31]: (2124, 220)
```

```
In [32]: train_df['Gene'].head()
```

```
Out[32]: 12      CBL  

192     EGFR  

2558    BRCA1  

1062    EWSR1  

2395     NF1  

Name: Gene, dtype: object
```

```
In [33]: gene_vectorizer.get_feature_names()
```

```
Out[33]: ['abl1',  

'acvr1',  

'ago2',  

'akt1',  

'akt2',  

'akt3',  

'alk',  

'apc',  

'ar',  

'araf',  

'arid1b',  

'arid2',  

'atm',  

'atrx',  

'aurka',  

'axl',  

'b2m',  

'bap1',  

'bard1',  

'bcl11a']
```



```
In [34]: # creating a pandas dataframe of the vectorized features
df_gene_train = pd.DataFrame(train_gene_feature_onehotCoding.toarray())
df_gene_test = pd.DataFrame(test_gene_feature_onehotCoding.toarray(), columns=
df_gene_cv = pd.DataFrame(cv_gene_feature_onehotCoding.toarray(), columns=
```

```
In [35]: print("train_gene_feature_onehotCoding is converted feature using one-hot encoding method. The shape of gene feature: (2124, 220)")

train_gene_feature_onehotCoding is converted feature using one-hot encoding method. The shape of gene feature: (2124, 220)
```

Q4. How good is this gene feature in predicting y_i ?

There are many ways to estimate how good a feature is, in predicting y_i . One of the good methods is to build a proper ML model using just this feature. In this case, we will build a logistic regression model using only Gene feature (one hot encoded) to predict y_i .

```
In [36]: alpha = [10 ** x for x in range(-5, 1)] # hyperparam for SGD classifier

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=None)
    clf.fit(train_gene_feature_onehotCoding, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_gene_feature_onehotCoding, y_train)
    predict_y = sig_clf.predict_proba(cv_gene_feature_onehotCoding)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
```

```


plt.xlabel("alpha")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
clf.fit(train_gene_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_gene_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log
predict_y = sig_clf.predict_proba(cv_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross val
predict_y = sig_clf.predict_proba(test_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log

```

For values of alpha = 1e-05 The log loss is: 1.3944084716286247
 For values of alpha = 0.0001 The log loss is: 1.231383844235684
 For values of alpha = 0.001 The log loss is: 1.2512813436315235
 For values of alpha = 0.01 The log loss is: 1.3731077811425252
 For values of alpha = 0.1 The log loss is: 1.4776878210761122
 For values of alpha = 1 The log loss is: 1.510967003941564



For values of best alpha = 0.0001 The train log loss is: 1.0461554141926563
 For values of best alpha = 0.0001 The cross validation log loss is: 1.231383844235684
 For values of best alpha = 0.0001 The test log loss is: 1.2098656591038013

Q5. Is the Gene feature stable across all the data sets (Test, Train, Cross validation)?

Ans. Yes, it is. Otherwise, the CV and Test errors would be significantly more than train error.

```
In [37]: print("Q6. How many data points in Test and CV datasets are covered by

test_coverage=test_df[test_df['Gene'].isin(list(set(train_df['Gene'])))
cv_coverage=cv_df[cv_df['Gene'].isin(list(set(train_df['Gene'])))].shape[0]

print('Ans\n1. In test data',test_coverage, 'out of',test_df.shape[0],
print('2. In cross validation data',cv_coverage, 'out of ',cv_df.shape[0])
```

Q6. How many data points in Test and CV datasets are covered by the 224 genes in train dataset?

Ans

1. In test data 645 out of 665 : 96.99248120300751

2. In cross validation data 507 out of 532 : 95.30075187969925

3.2.2 Univariate Analysis on Variation Feature

Q7. Variation, What type of feature is it ?

Ans. Variation is a categorical variable

Q8. How many categories are there?

```
In [38]: unique_variations = train_df['Variation'].value_counts()
print('Number of Unique Variations :', unique_variations.shape[0])
# the top 10 variations that occurred most
print(unique_variations.head(10))
```

Number of Unique Variations : 1930

Truncating_Mutations 57

Deletion 43

Amplification 42

Fusions 25

T58I 3

Overexpression 3

G12V 3

R170W 2

Q61L 2


Y64A 2

Name: Variation, dtype: int64

```
In [39]: print("Ans: There are", unique_variations.shape[0] , "different categor.
```


Ans: There are 1930 different categories of variations in the train data, and they are distributed as follows

```
In [40]: s = sum(unique_variations.values);  
h = unique_variations.values/s;  
plt.plot(h, label="Histogram of Variations")  
plt.xlabel('Index of a Variation')  
plt.ylabel('Number of Occurances')  
plt.legend()  
plt.grid()  
plt.show()
```



```
In [41]: c = np.cumsum(h)
print(c)
plt.plot(c,label='Cumulative distribution of Variations')
plt.grid()
plt.legend()
plt.show()

[0.02683616 0.04708098 0.06685499 ... 0.99905838 0.99952919 1.
]
```



Q9. How to featurize this Variation feature ?

Ans. There are two ways we can featurize this variable check out this video:

<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>
[\(https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/\)](https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/)

1. One hot Encoding
2. Response coding

We will be using both these methods to featurize the Variation Feature

```
In [42]: # alpha is used for laplace smoothing
alpha = 1
# train gene feature
train_variation_feature_responseCoding = np.array(get_gv_feature(alpha
# test gene feature
test_variation_feature_responseCoding = np.array(get_gv_feature(alpha,
# cross validation gene feature
cv_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "
```

```
In [43]: print("train_variation_feature_responseCoding is a converted feature using the response coding method. The shape of Variation feature: (2124, 9)")
```

```
In [44]: # one-hot encoding of variation feature.
variation_vectorizer = TfidfVectorizer(max_features=1700)
train_variation_feature_onehotCoding = variation_vectorizer.fit_transform(train_variation_feature_responseCoding)
test_variation_feature_onehotCoding = variation_vectorizer.transform(test_variation_feature_responseCoding)
cv_variation_feature_onehotCoding = variation_vectorizer.transform(cv_variation_feature_responseCoding)
```

```
In [45]: df_var_train = pd.DataFrame(train_variation_feature_onehotCoding.toarray(), columns=train_variation_feature_onehotCoding.columns)
df_var_test = pd.DataFrame(test_variation_feature_onehotCoding.toarray(), columns=test_variation_feature_onehotCoding.columns)
df_var_cv = pd.DataFrame(cv_variation_feature_onehotCoding.toarray(), columns=cv_variation_feature_onehotCoding.columns)
```

```
In [46]: print("train_variation_feature_onehotEncoded is converted feature using the one-hot encoding method. The shape of Variation feature: (2124, 1700)")
```

Q10. How good is this Variation feature in predicting y_i?

Let's build a model just like the earlier!

```
In [47]: alpha = [10 ** x for x in range(-5, 1)]

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=None)
    clf.fit(train_variation_feature_onehotCoding, y_train)
```

```

sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_variation_feature_onehotCoding, y_train)
predict_y = sig_clf.predict_proba(cv_variation_feature_onehotCoding)

cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_))
print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
                    class_weight='balanced')
clf.fit(train_variation_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_variation_feature_onehotCoding, y_train)


predict_y = sig_clf.predict_proba(train_variation_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is: ", log_loss(y_train, predict_y, labels=clf.classes_))
predict_y = sig_clf.predict_proba(cv_variation_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is: ", log_loss(y_cv, predict_y, labels=clf.classes_))
predict_y = sig_clf.predict_proba(test_variation_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is: ", log_loss(y_test, predict_y, labels=clf.classes_))

```

```

For values of alpha = 1e-05 The log loss is: 1.7207189013959114
For values of alpha = 0.0001 The log loss is: 1.6997689084309093
For values of alpha = 0.001 The log loss is: 1.7026016503896444
For values of alpha = 0.01 The log loss is: 1.7131165722610504
For values of alpha = 0.1 The log loss is: 1.7173932048592393
For values of alpha = 1 The log loss is: 1.7186310786492505

```



Alpha 1.5

For values of best alpha = 0.0001 The train log loss is: 0.86508446
 86325048
 For values of best alpha = 0.0001 The cross validation log loss is:
 1.6997689084309093
 For values of best alpha = 0.0001 The test log loss is: 1.703251480
 2236773

Q11. Is the Variation feature stable across all the data sets (Test, Train, Cross validation)?

Ans. Not sure! But lets be very sure using the below analysis.

```
In [48]: print("Q12. How many data points are covered by total ", unique_variat
test_coverage=test_df[test_df['Variation'].isin(list(set(train_df['Var
cv_coverage=cv_df[cv_df['Variation'].isin(list(set(train_df['Variation
print('Ans\n1. In test data',test_coverage, 'out of',test_df.shape[0],
print('2. In cross validation data',cv_coverage, 'out of ',cv_df.shape
```

Q12. How many data points are covered by total 1930 genes in test and cross validation data sets?

Ans

1. In test data 69 out of 665 : 10.37593984962406
2. In cross validation data 55 out of 532 : 10.338345864661653

3.2.3 Univariate Analysis on Text Feature

1. How many unique words are present in train data?
2. How are word frequencies distributed?
3. How to featurize text field?
4. Is the text feature useful in prediciting y_i?
5. Is the text feature stable across train, test and CV datasets?

```
In [49]: # cls_text is a data frame
# for every row in data fram consider the 'TEXT'
# split the words by space
# make a dict with those words
# increment its count whenever we see that word

def extract_dictionary_paddle(cls_text):
    dictionary = defaultdict(int)
    for index, row in cls_text.iterrows():
        for word in row['TEXT'].split():
            dictionary[word] +=1
    return dictionary
```



```
In [50]: import math
#https://stackoverflow.com/a/1602964
def get_text_responseCoding(df):
    text_feature_responseCoding = np.zeros((df.shape[0],9))
    for i in range(0,9):
        row_index = 0
        for index, row in df.iterrows():
            sum_prob = 0
            for word in row['TEXT'].split():
                sum_prob += math.log(((dict_list[i].get(word,0)+10 ))/(
            text_feature_responseCoding[row_index][i] = math.exp(sum_p
            row_index += 1
    return text_feature_responseCoding
```

```
In [51]: # building a CountVectorizer with all the words that occurred minimum 3
text_vectorizer = TfidfVectorizer(min_df=5,ngram_range=(1,2),max_featu
train_text_feature_onehotCoding = text_vectorizer.fit_transform(train_
# getting all the feature names (words)
train_text_features= text_vectorizer.get_feature_names()

# train_text_feature_onehotCoding.sum(axis=0).A1 will sum every row and
train_text_fea_counts = train_text_feature_onehotCoding.sum(axis=0).A1

# zip(list(text_features),text_fea_counts) will zip a word with its num
text_fea_dict = dict(zip(list(train_text_features),train_text_fea_count

print("Total number of unique words in train data :", len(train_text_f
```

Total number of unique words in train data : 20000

```
In [52]: dict_list = []
# dict_list =[] contains 9 dictionaries each corresponds to a class
for i in range(1,10):
    cls_text = train_df[train_df['Class']==i]
    # build a word dict based on the words in that class
    dict_list.append(extract_dictionary_paddle(cls_text))
    # append it to dict_list

# dict_list[i] is build on i'th class text data
# total_dict is build on whole training text data
total_dict = extract_dictionary_paddle(train_df)

confuse_array = []
for i in train_text_features:
    ratios = []
    max_val = -1
    for j in range(0,9):
        ratios.append((dict_list[j][i]+10)/(total_dict[i]+90))
    confuse_array.append(ratios)
confuse_array = np.array(confuse_array)
```

```
In [53]: #response coding of text features
train_text_feature_responseCoding = get_text_responsecoding(train_df)
test_text_feature_responseCoding = get_text_responsecoding(test_df)
cv_text_feature_responseCoding = get_text_responsecoding(cv_df)
```

```
In [54]: # https://stackoverflow.com/a/16202486
# we convert each row values such that they sum to 1
train_text_feature_responseCoding = (train_text_feature_responseCoding.T+1).T
test_text_feature_responseCoding = (test_text_feature_responseCoding.T+1).T
cv_text_feature_responseCoding = (cv_text_feature_responseCoding.T+1).T
```

```
In [55]: # don't forget to normalize every feature
train_text_feature_onehotCoding = normalize(train_text_feature_onehotCoding)

# we use the same vectorizer that was trained on train data
test_text_feature_onehotCoding = text_vectorizer.transform(test_df['TEXT'])
# don't forget to normalize every feature
test_text_feature_onehotCoding = normalize(test_text_feature_onehotCoding)

# we use the same vectorizer that was trained on train data
cv_text_feature_onehotCoding = text_vectorizer.transform(cv_df['TEXT'])
# don't forget to normalize every feature
cv_text_feature_onehotCoding = normalize(cv_text_feature_onehotCoding)
```

```
In [56]: #https://stackoverflow.com/a/2258273/4084039
sorted_text_fea_dict = dict(sorted(text_fea_dict.items(), key=lambda x: x[1]))
sorted_text_occur = np.array(list(sorted_text_fea_dict.values()))
```

```
In [57]: # Number of words for a given frequency.
print(Counter(sorted_text_occur))
```

```
Counter({0.5928400952859082: 39, 0.7544202940499756: 32, 0.666516722
1822955: 26, 0.6388896632493938: 20, 0.8886889629097272: 16, 1.06481
61054156562: 15, 0.9085949753176206: 14, 1.0389694259244064: 10, 3.1
937793216260233: 8, 0.851852884332525: 8, 0.7415078640131119: 7, 0.6
975390741666331: 7, 0.889260142928862: 6, 0.6641917020380229: 6, 0.5
572043451611401: 6, 2.589247760039847: 5, 2.1296322108313124: 5, 1.4
907425475819174: 5, 1.0836875989439445: 5, 0.9421724920023471: 5, 0.
6171172984548712: 5, 0.5800558825457663: 5, 4.487153349100253: 4, 2.
422919934180324: 4, 1.621224820902286: 4, 1.514324958862703: 4, 1.50
88405880999511: 4, 1.333033444364591: 4, 1.2982056046452644: 4, 1.21
1459967090162: 4, 1.1108612036371595: 4, 0.753737993601878: 4, 0.694
994699185837: 4, 0.6686452141933684: 4, 0.5255716752962893: 4, 0.522
4245169756063: 4, 0.4785805665793877: 4, 2.7685218740807045: 3, 2.63
46823125623207: 3, 2.342595431914442: 3, 1.8174878465033697: 3, 1.69
59104856042255: 3, 1.555205685092022: 3, 1.2777793264987876: 3, 1.13
16304410749625: 3, 1.0294102356167: 3, 0.91333191791486: 3, 0.817018
3652277363: 3, 0.7378770862301951: 3, 0.7273698269507515: 3, 0.72167
13210917822: 3, 0.6980113222458281: 3, 0.6491028023226322: 3, 0.6123
100131839002: 3, 0.5988436571197083: 3, 0.5559957593486697: 3, 0.532
...})
```

```
In [58]: # Train a Logistic regression+Calibration model using text features wh
alpha = [10 ** x for x in range(-5, 1)]

# read more about SGDClassifier() at http://scikit-learn.org/stable/mo
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.1
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, l
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stoc
# predict(X) Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_stat
    clf.fit(train_text_feature_onehotCoding, y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_text_feature_onehotCoding, y_train)
    predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.clas
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv
```

```

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_log_error_arr
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
clf.fit(train_text_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_text_feature_onehotCoding, y_train)


predict_y = sig_clf.predict_proba(train_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log
predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross val
predict_y = sig_clf.predict_proba(test_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log

```

```

For values of alpha = 1e-05 The log loss is: 1.149056113623601
For values of alpha = 0.0001 The log loss is: 1.122895305061825
For values of alpha = 0.001 The log loss is: 1.152449660058874
For values of alpha = 0.01 The log loss is: 1.2802353839118785
For values of alpha = 0.1 The log loss is: 1.4137492341435582
For values of alpha = 1 The log loss is: 1.605535246074835

```



```

For values of best alpha = 0.0001 The train log loss is: 0.86219088
58537646
For values of best alpha = 0.0001 The cross validation log loss is:
1.122895305061825
For values of best alpha = 0.0001 The test log loss is: 1.131701569

```

2145014

```
In [59]: df_text_train = pd.DataFrame(train_text_feature_onehotCoding.toarray())
df_text_test = pd.DataFrame(test_text_feature_onehotCoding.toarray()),
df_text_cv = pd.DataFrame(cv_text_feature_onehotCoding.toarray(), colum
```

Q. Is the Text feature stable across all the data sets (Test, Train, Cross validation)?

Ans. Yes, it seems like!

```
In [60]: def get_intersec_text(df):
df_text_vec = TfidfVectorizer(min_df=3)
df_text_fea = df_text_vec.fit_transform(df['TEXT'])
df_text_features = df_text_vec.get_feature_names()

df_text_fea_counts = df_text_fea.sum(axis=0).A1
df_text_fea_dict = dict(zip(list(df_text_features),df_text_fea_counts))
len1 = len(set(df_text_features))
len2 = len(set(train_text_features) & set(df_text_features))
return len1,len2
```

```
In [61]: len1,len2 = get_intersec_text(test_df)
print(np.round((len2/len1)*100, 3), "% of word of test data appeared in train data")
len1,len2 = get_intersec_text(cv_df)
print(np.round((len2/len1)*100, 3), "% of word of Cross Validation appeared in train data")
```

```
30.223 % of word of test data appeared in train data
35.099 % of word of Cross Validation appeared in train data
```

Univariate analysis of no_of_words

```
In [62]: alpha = [10 ** x for x in range(-5, 1)]

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=0)
    clf.fit(train_df["no_of_words"].values.reshape(-1,1), y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_df.no_of_words.values.reshape(-1,1), y_train)
    predict_y = sig_clf.predict_proba(cv_df.no_of_words.values.reshape(-1,1))

    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_))
    print('For values of alpha = ', i, "The log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='r')
```

```

for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
clf.fit(train_df.no_of_words.values.reshape(-1,1), y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_df.no_of_words.values.reshape(-1,1), y_train)


predict_y = sig_clf.predict_proba(train_df.no_of_words.values.reshape(-1,1))
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is: ", train_log_loss)
predict_y = sig_clf.predict_proba(cv_df.no_of_words.values.reshape(-1,1))
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is: ", cv_log_loss)
predict_y = sig_clf.predict_proba(test_df.no_of_words.values.reshape(-1,1))
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is: ", test_log_loss)

```

```

For values of alpha = 1e-05 The log loss is: 1.828456718086271
For values of alpha = 0.0001 The log loss is: 1.8284567180762197
For values of alpha = 0.001 The log loss is: 1.8280692399388123
For values of alpha = 0.01 The log loss is: 1.8276398372891882
For values of alpha = 0.1 The log loss is: 1.8285563158071414
For values of alpha = 1 The log loss is: 1.8292590645707076

```



```

For values of best alpha = 0.01 The train log loss is: 1.8123028601858309
For values of best alpha = 0.01 The cross validation log loss is: 1.8276398372891882
For values of best alpha = 0.01 The test log loss is: 1.8112376518482927

```

Univariate analysis of no_of_characters

```
In [63]: alpha = [10 ** x for x in range(-5, 1)]

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=0)
    clf.fit(train_df["no_of_characters"].values.reshape(-1,1), y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_df.no_of_characters.values.reshape(-1,1), y_train)
    predict_y = sig_clf.predict_proba(cv_df.no_of_characters.values.reshape(-1,1))


    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_))


fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=0)
clf.fit(train_df.no_of_characters.values.reshape(-1,1), y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_df.no_of_characters.values.reshape(-1,1), y_train)

predict_y = sig_clf.predict_proba(train_df.no_of_characters.values.reshape(-1,1))
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_))
predict_y = sig_clf.predict_proba(cv_df.no_of_characters.values.reshape(-1,1))
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_))
predict_y = sig_clf.predict_proba(test_df.no_of_characters.values.reshape(-1,1))
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_))
```

```
For values of alpha = 1e-05 The log loss is: 1.828162256019388
For values of alpha = 0.0001 The log loss is: 1.8281622560192532
For values of alpha = 0.001 The log loss is: 1.8281622559915802
For values of alpha = 0.01 The log loss is: 1.8281622560558215
For values of alpha = 0.1 The log loss is: 1.828236977767708
For values of alpha = 1 The log loss is: 1.8270168837018603
```





For values of best alpha = 1 The train log loss is: 1.808939434850306

For values of best alpha = 1 The cross validation log loss is: 1.8270168837018603

For values of best alpha = 1 The test log loss is: 1.8082689185107592

Univariate analysis of gene_variation

```
In [64]: gene_var = train_df['gene_variation'].value_counts()
print('Count of Gene and Variation which are unique :', gene_var.shape)
# the top 10 variations that occurred most
print("Head", gene_var.head(10))
print("Tail", gene_var.head(10))
```

Count of Gene and Variation which are unique : 2124

Head SMAD3 S425C 1

TET2 R1896M 1

MTOR D2512G 1

CTNNB1 G34E 1

KIT D816E 1

TSC2 R611W 1

RASA1 Y472H 1

ALK D1349H 1

ALK S1206R 1

IDH2 R172K 1

Name: gene_variation, dtype: int64

Tail SMAD3 S425C 1

TET2 R1896M 1

MTOR D2512G 1

CTNNB1 G34E 1

KIT D816E 1

TSC2 R611W 1

RASA1 Y472H 1

ALK D1349H 1

ALK S1206R 1

IDH2 R172K 1

Name: gene_variation, dtype: int64


```
In [65]: # Featurizing the Gene_and_Variation Feature
# alpha is used for laplace smoothing
alpha = 1
# train gene feature
train_gene_and_variation_feature_responseCoding = np.array(get_gv_featu
# test gene feature
test_gene_and_variation_feature_responseCoding = np.array(get_gv_featu
# cross validation gene feature
cv_gene_and_variation_feature_responseCoding = np.array(get_gv_feature
```

```
In [66]: # one-hot encoding of gene_and_variation feature.
gene_variation_vectorizer = TfidfVectorizer()
train_gene_and_variation_feature_onehotCoding = gene_variation_vectorizer
test_gene_and_variation_feature_onehotCoding = gene_variation_vectorizer
cv_gene_and_variation_feature_onehotCoding = gene_variation_vectorizer
```

```
In [67]: df_geneandvar_train = pd.DataFrame(train_gene_and_variation_feature_one
df_geneandvar_test = pd.DataFrame(test_gene_and_variation_feature_onehot
df_geneandvar_cv = pd.DataFrame(cv_gene_and_variation_feature_onehotCo
```

```
In [68]: alpha = [10 ** x for x in range(-5, 1)]

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=0)
    clf.fit(df_geneandvar_train, y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(df_geneandvar_train, y_train)
    predict_y = sig_clf.predict_proba(df_geneandvar_cv)

    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=0)
clf.fit(df_geneandvar_train, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(df_geneandvar_train, y_train)
```

```


predict_y = sig_clf.predict_proba(df_geneandvar_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log
predict_y = sig_clf.predict_proba(df_geneandvar_cv)
print('For values of best alpha = ', alpha[best_alpha], "The cross val
predict_y = sig_clf.predict_proba(df_geneandvar_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log

```

```

For values of alpha = 1e-05 The log loss is: 1.1788850909646191
For values of alpha = 0.0001 The log loss is: 1.1564038096282738
For values of alpha = 0.001 The log loss is: 1.1666389694783292
For values of alpha = 0.01 The log loss is: 1.2004413194107715
For values of alpha = 0.1 The log loss is: 1.2037014208637031
For values of alpha = 1 The log loss is: 1.2424739088579682

```



```

For values of best alpha = 0.0001 The train log loss is: 0.50735322
79878828
For values of best alpha = 0.0001 The cross validation log loss is:
1.1564038096282738
For values of best alpha = 0.0001 The test log loss is: 1.113590880
0719008

```

Univariate analysis of word_count_of_5k

```

In [69]: alpha = [10 ** x for x in range(-5, 1)]

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=0)
    clf.fit(train_df["word_count_of_5k"].values.reshape(-1,1), y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_df.word_count_of_5k.values.reshape(-1,1), y_train)
    predict_y = sig_clf.predict_proba(cv_df.word_count_of_5k.values.re

```

```

cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_))
print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_))


fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
                    cv_loss_function=log_loss, nb_epochs=10000)
clf.fit(train_df.word_count_of_5k.values.reshape(-1, 1), y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_df.word_count_of_5k.values.reshape(-1, 1), y_train)

predict_y = sig_clf.predict_proba(train_df.word_count_of_5k.values.reshape(-1, 1))
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is: ", log_loss(y_train, predict_y, labels=clf.classes_))
predict_y = sig_clf.predict_proba(cv_df.word_count_of_5k.values.reshape(-1, 1))
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is: ", log_loss(y_cv, predict_y, labels=clf.classes_))
predict_y = sig_clf.predict_proba(test_df.word_count_of_5k.values.reshape(-1, 1))
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is: ", log_loss(y_test, predict_y, labels=clf.classes_))

```

For values of alpha = 1e-05 The log loss is: 1.8356817613570136
 For values of alpha = 0.0001 The log loss is: 1.8356817615055787
 For values of alpha = 0.001 The log loss is: 1.8356817594697372
 For values of alpha = 0.01 The log loss is: 1.8356817656528432
 For values of alpha = 0.1 The log loss is: 1.8356817452536163
 For values of alpha = 1 The log loss is: 1.8356817624766348



For values of best alpha = 0.1 The train log loss is: 1.8182613600460846
 For values of best alpha = 0.1 The cross validation log loss is: 1.8356817452536163

For values of best alpha = 0.1 The test log loss is: 1.8127482506828

Univariate analysis of character_count_of_50k

```
In [70]: alpha = [10 ** x for x in range(-5, 1)]

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=0)
    clf.fit(train_df["character_count_of_50k"].values.reshape(-1,1), y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_df.character_count_of_50k.values.reshape(-1,1), y_train)
    predict_y = sig_clf.predict_proba(cv_df.character_count_of_50k.values.reshape(-1,1))


    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_))


fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=0)
clf.fit(train_df.character_count_of_50k.values.reshape(-1,1), y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_df.character_count_of_50k.values.reshape(-1,1), y_train)

predict_y = sig_clf.predict_proba(train_df.character_count_of_50k.values.reshape(-1,1))
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is: ", log_loss(y_train, predict_y, labels=clf.classes_))
predict_y = sig_clf.predict_proba(cv_df.character_count_of_50k.values.reshape(-1,1))
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is: ", log_loss(y_cv, predict_y, labels=clf.classes_))
predict_y = sig_clf.predict_proba(test_df.character_count_of_50k.values.reshape(-1,1))
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is: ", log_loss(y_test, predict_y, labels=clf.classes_))
```

For values of alpha = 1e-05 The log loss is: 1.8271522109805236
 For values of alpha = 0.0001 The log loss is: 1.8271522106894227
 For values of alpha = 0.001 The log loss is: 1.8271522119019332
 For values of alpha = 0.01 The log loss is: 1.8271522155237467
 For values of alpha = 0.1 The log loss is: 1.8271522105135294
 For values of alpha = 1 The log loss is: 1.8271522094197878





For values of best alpha = 1 The train log loss is: 1.8105178137952094

For values of best alpha = 1 The cross validation log loss is: 1.8271522094197878

For values of best alpha = 1 The test log loss is: 1.8082753115230226

4. Machine Learning Models

In [71]: *#Data preparation for ML models.*

#Misc. functionns for ML models

```
def predict_and_plot_confusion_matrix(train_x, train_y, test_x, test_y,
                                       clf.fit(train_x, train_y)
                                       sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
                                       sig_clf.fit(train_x, train_y)
                                       pred_y = sig_clf.predict(test_x)

   # for calculating log_loss we willl provide the array of probabili
   print("Log loss :", log_loss(test_y, sig_clf.predict_proba(test_x))
   # calculating the number of data points that are misclassified
   print("Number of mis-classified points :", np.count_nonzero((pred_
   plot_confusion_matrix(test_y, pred_y)
```

In [72]: **def** report_log_loss(train_x, train_y, test_x, test_y, clf):

```
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    sig_clf_probs = sig_clf.predict_proba(test_x)
    return log_loss(test_y, sig_clf_probs, eps=1e-15)
```

```

In [110]: # this function will be used just for naive bayes
# for the given indices, we will print the name of the features
# and we will check whether the feature present in the test point text
def get_impfeature_names(indices, text, gene, var, no_features):
    gene_count_vec = TfidfVectorizer()
    var_count_vec = TfidfVectorizer()
    text_count_vec = TfidfVectorizer(min_df=3,max_features=1000)

    gene_vec = gene_count_vec.fit(train_df['Gene'])
    var_vec = var_count_vec.fit(train_df['Variation'])
    text_vec = text_count_vec.fit(train_df['TEXT'])

    fea1_len = len(gene_vec.get_feature_names())
    fea2_len = len(var_count_vec.get_feature_names())

    word_present = 0
    for i,v in enumerate(indices):
        if (v < fea1_len):
            word = gene_vec.get_feature_names()[v]
            yes_no = True if word == gene else False
            if yes_no:
                word_present += 1
                print(i, "Gene feature [{}] present in test data point".format(word))
        elif (v < fea1_len+fea2_len):
            word = var_vec.get_feature_names()[v-(fea1_len)]
            yes_no = True if word == var else False
            if yes_no:
                word_present += 1
                print(i, "variation feature [{}] present in test data point".format(word))
        else:
            word = text_vec.get_feature_names()[v-(fea1_len+fea2_len)]
            yes_no = True if word in text.split() else False
            if yes_no:
                word_present += 1
                print(i, "Text feature [{}] present in test data point".format(word))

    print("Out of the top ",no_features," features ", word_present, " are present")

```

Stacking the three types of features

```
In [74]: #target variables
train_y = train_df['Class'].values
test_y = test_df['Class'].values
cv_y = cv_df['Class'].values

# concatenating all the vectorized dataframes
df_gene_var_train = pd.concat([df_gene_train, df_var_train], axis=1)
df_gene_var_test = pd.concat([df_gene_test, df_var_test], axis=1)
df_gene_var_cv = pd.concat([df_gene_cv, df_var_cv], axis=1)

df_gene_and_var_train = pd.concat([df_gene_var_train, df_geneandvar_train], axis=1)
df_gene_and_var_test = pd.concat([df_gene_var_test, df_geneandvar_test], axis=1)
df_gene_and_var_cv = pd.concat([df_gene_var_cv, df_geneandvar_cv], axis=1)

df_train = pd.concat([df_gene_and_var_train, df_text_train], axis=1)
df_test = pd.concat([df_gene_and_var_test, df_text_test], axis=1)
df_cv = pd.concat([df_gene_and_var_cv, df_text_cv], axis=1)
```

```

In [75]: # scaling the text_count feature
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
train_df["no_of_words"] = scaler.fit_transform(train_df["no_of_words"])
test_df["no_of_words"] = scaler.fit_transform(test_df["no_of_words"].values)
cv_df["no_of_words"] = scaler.fit_transform(cv_df["no_of_words"].values)

train_df["no_of_characters"] = scaler.fit_transform(train_df["no_of_characters"])
test_df["no_of_characters"] = scaler.fit_transform(test_df["no_of_characters"].values)
cv_df["no_of_characters"] = scaler.fit_transform(cv_df["no_of_characters"].values)

train_df["word_count_of_5k"] = scaler.fit_transform(train_df["word_count_of_5k"])
test_df["word_count_of_5k"] = scaler.fit_transform(test_df["word_count_of_5k"].values)
cv_df["word_count_of_5k"] = scaler.fit_transform(cv_df["word_count_of_5k"].values)

train_df["character_count_of_50k"] = scaler.fit_transform(train_df["character_count_of_50k"])
test_df["character_count_of_50k"] = scaler.fit_transform(test_df["character_count_of_50k"].values)
cv_df["character_count_of_50k"] = scaler.fit_transform(cv_df["character_count_of_50k"].values)

df_train["no_of_words"] = train_df.no_of_words.values
df_train["no_of_characters"] = train_df.no_of_characters.values
df_train["word_count_of_5k"] = train_df.word_count_of_5k.values
df_train["character_count_of_50k"] = train_df.character_count_of_50k.values

df_test["no_of_words"] = test_df.no_of_words.values
df_test["no_of_characters"] = test_df.no_of_characters.values
df_test["word_count_of_5k"] = test_df.word_count_of_5k.values
df_test["character_count_of_50k"] = test_df.character_count_of_50k.values

df_cv["no_of_words"] = cv_df.no_of_words.values
df_cv["no_of_characters"] = cv_df.no_of_characters.values
df_cv["word_count_of_5k"] = cv_df.word_count_of_5k.values
df_cv["character_count_of_50k"] = cv_df.character_count_of_50k.values

```



```

In [76]: train_gene_var_responseCoding = np.hstack((train_gene_feature_responseCoding, train_gene_var_responseCoding))
test_gene_var_responseCoding = np.hstack((test_gene_feature_responseCoding, test_gene_var_responseCoding))
cv_gene_var_responseCoding = np.hstack((cv_gene_feature_responseCoding, cv_gene_var_responseCoding))

train_geneandvar_responseCoding = np.hstack((train_gene_var_responseCoding, train_geneandvar_responseCoding))
test_geneandvar_responseCoding = np.hstack((test_gene_var_responseCoding, test_geneandvar_responseCoding))
cv_geneandvar_responseCoding = np.hstack((cv_gene_var_responseCoding, cv_geneandvar_responseCoding))

train_x_responseCoding = np.hstack((train_geneandvar_responseCoding, train_x_responseCoding))
test_x_responseCoding = np.hstack((test_geneandvar_responseCoding, test_x_responseCoding))
cv_x_responseCoding = np.hstack((cv_geneandvar_responseCoding, cv_x_responseCoding))

train_no_of_words = np.column_stack((train_x_responseCoding, train_df.no_of_words))
test_no_of_words = np.column_stack((test_x_responseCoding, test_df.no_of_words))
cv_no_of_words = np.column_stack((cv_x_responseCoding, cv_df.no_of_words))

train_no_of_characters = np.column_stack((train_no_of_words, train_df.no_of_characters))
test_no_of_characters = np.column_stack((test_no_of_words, test_df.no_of_characters))
cv_no_of_characters = np.column_stack((cv_no_of_words, cv_df.no_of_characters))

train_word_count_of_5k = np.column_stack((train_no_of_characters, train_df.word_count_of_5k))
test_word_count_of_5k = np.column_stack((test_no_of_characters, test_df.word_count_of_5k))
cv_word_count_of_5k = np.column_stack((cv_no_of_characters, cv_df.word_count_of_5k))

train_x_response = np.column_stack((train_word_count_of_5k, train_df.character_count_of_5k))
test_x_response = np.column_stack((test_word_count_of_5k, test_df.character_count_of_5k))
cv_x_response = np.column_stack((cv_word_count_of_5k, cv_df.character_count_of_5k))

train_x_onehotCoding = df_train
test_x_onehotCoding = df_test
cv_x_onehotCoding = df_cv

train_x_responseCoding = pd.DataFrame(train_x_response)
test_x_responseCoding = pd.DataFrame(test_x_response)
cv_x_responseCoding = pd.DataFrame(cv_x_response)

```

```

In [77]: print("One hot encoding features :")
print("(number of data points * number of features) in train data = ",
print("(number of data points * number of features) in test data = ",
print("(number of data points * number of features) in cross validation data = ")

```

```

One hot encoding features :
(number of data points * number of features) in train data = (2124, 24071)
(number of data points * number of features) in test data = (665, 24071)
(number of data points * number of features) in cross validation data = (532, 24071)

```

```
In [78]: print(" Response encoding features :")
print("(number of data points * number of features) in train data = ",
print("(number of data points * number of features) in test data = ",
print("(number of data points * number of features) in cross validation data = ")

Response encoding features :
(number of data points * number of features) in train data = (2124, 40)
(number of data points * number of features) in test data = (665, 40)
(number of data points * number of features) in cross validation data = (532, 40)
```

4.1. Base Line Model

4.1.1. Naive Bayes

4.1.1.1. Hyper parameter tuning

```
In [79]: # find more about Multinomial Naive base function here http://scikit-learn.org/stable/modules/generated/sklearn.naive\_bayes.MultinomialNB.html
# -----
# default parameters
# sklearn.naive_bayes.MultinomialNB(alpha=1.0, fit_prior=True, class_prior=None)

# some of methods of MultinomialNB()
# fit(X, y[, sample_weight]) Fit Naive Bayes classifier according to training set (X, y)
# predict(X) Perform classification on an array of test vectors X.
# predict_log_proba(X) Return log-probability estimates for the test vectors X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default parameters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=5)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
# -----

alpha = [0.00001, 0.0001, 0.001, 0.1, 1, 10, 100, 1000]
```

```

alpha = [0.00001, 0.00001, 0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = MultinomialNB(alpha=i)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_))
    # to avoid rounding error while multiplying probabilities we use log
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(np.log10(alpha), cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (np.log10(alpha[i]), cv_log_error_array[i]))
plt.grid()
plt.xticks(np.log10(alpha))
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = MultinomialNB(alpha=alpha[best_alpha])
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is: ", log_loss(train_y, predict_y, labels=clf.classes_))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is: ", log_loss(cv_y, predict_y, labels=clf.classes_))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is: ", log_loss(test_y, predict_y, labels=clf.classes_))


```

```

for alpha = 1e-05
Log Loss : 1.2128577780836407
for alpha = 0.0001
Log Loss : 1.21817052556479
for alpha = 0.001
Log Loss : 1.2158982215448673
for alpha = 0.1
Log Loss : 1.2477365646690524
for alpha = 1
Log Loss : 1.2890814072732695
for alpha = 10
Log Loss : 1.4216474071762748
for alpha = 100
Log Loss : 1.4136778012386413

```

```
for alpha = 1000
Log Loss : 1.424206335067884
```



For values of best alpha = 1e-05 The train log loss is: 0.7945349704848299

For values of best alpha = 1e-05 The cross validation log loss is: 1.2128577780836407

For values of best alpha = 1e-05 The test log loss is: 1.2321949417796758

4.1.1.2. Testing the model with best hyper paramters

```
In [80]: # find more about Multinomial Naive base function here http://scikit-learn.org/
# -----
# default paramters
# sklearn.naive_bayes.MultinomialNB(alpha=1.0, fit_prior=True, class_prior=None)

# some of methods of MultinomialNB()
# fit(X, y[, sample_weight]) Fit Naive Bayes classifier according to training set X
# predict(X) Perform classification on an array of test vectors X.
# predict_log_proba(X) Return log-probability estimates for the test vectors X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-2019
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid')
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Return probability estimates for the test vectors X.
```


```
# predict_proba(x) POSTERIOR PROBABILITIES OF CLASSIFICATION
# -----

clf = MultinomialNB(alpha=alpha[best_alpha])
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)
sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
# to avoid rounding error while multiplying probabilities we use log-prob
print("Log Loss :", log_loss(cv_y, sig_clf_probs))
print("Number of missclassified point :", np.count_nonzero((sig_clf.predict(cv_x_onehotCoding) != cv_y)))
plot_confusion_matrix(cv_y, sig_clf.predict(cv_x_onehotCoding))
```

Log Loss : 1.2128577780836407

Number of missclassified point : 0.35526315789473684


----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----





4.1.1.3. Feature Importance, Correctly classified point

```
In [81]: test_point_index = 12
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding.iloc[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotCoding.iloc[test_point_index])[0], 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_[predicted_cls-1][:, :no_feature])
print("-"*50)
```

```
Predicted Class : 7
Predicted Class Probabilities: [[0.0789 0.2437 0.0122 0.1006 0.0426
0.0374 0.4723 0.006 0.0062]]
Actual Class : 7
-----
```

4.1.1.4. Feature Importance, Incorrectly classified point

```
In [82]: test_point_index = 1
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding.iloc[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotCoding.iloc[test_point_index])[0], 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_[predicted_cls-1][:, :no_feature])
print("-"*50)
```

```
Predicted Class : 2
Predicted Class Probabilities: [[0.09 0.567 0.0135 0.1138 0.048
0.0422 0.1118 0.0068 0.007 ]]
Actual Class : 1
-----
```

4.2. K Nearest Neighbour Classification

4.2.1. Hyper parameter tuning

```
In [111]: # find more about KNeighborsClassifier() here http://scikit-learn.org/
# -----
# default parameter
# KNeighborsClassifier(n_neighbors=5, weights='uniform', algorithm='au
# metric='minkowski', metric_params=None, n_jobs=1, **kwargs)

# methods of
# fit(X, y) : Fit the model using X as training data and y as target v
# predict(X):Predict the class labels for the provided data
# predict_proba(X):Return probability estimates for the test data X.
#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.o
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, meth
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [5, 11, 15, 21, 31, 41, 51, 99]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = KNeighborsClassifier(n_neighbors=i)
    clf.fit(train_x_responseCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_responseCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf
    # to avoid rounding error while multiplying probabilitites we use lo
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
```

```

plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)


predict_y = sig_clf.predict_proba(train_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log")
predict_y = sig_clf.predict_proba(cv_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross val")
predict_y = sig_clf.predict_proba(test_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log")

```

```

for alpha = 5
Log Loss : 1.0985511326771829
for alpha = 11
Log Loss : 1.0732328472930994
for alpha = 15
Log Loss : 1.076528634036376
for alpha = 21
Log Loss : 1.0947602662954434
for alpha = 31
Log Loss : 1.1308694422401118
for alpha = 41
Log Loss : 1.156478238520508
for alpha = 51
Log Loss : 1.171171764652043
for alpha = 99
Log Loss : 1.2183618899099915

```



For values of best alpha = 11 The train log loss is: 0.7571029197349849

For values of best alpha = 11 The cross validation log loss is: 1.0732328472930994
 For values of best alpha = 11 The test log loss is: 1.0983942729743246

4.2.2. Testing the model with best hyper paramters


```
In [112]: # find more about KNeighborsClassifier() here http://scikit-learn.org/
# -----
# default parameter
# KNeighborsClassifier(n_neighbors=5, weights='uniform', algorithm='auto',
# metric='minkowski', metric_params=None, n_jobs=1, **kwargs)

# methods of
# fit(X, y) : Fit the model using X as training data and y as target variable
# predict(X):Predict the class labels for the provided data
# predict_proba(X):Return probability estimates for the test data X.
#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
#-----
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
predict_and_plot_confusion_matrix(train_x_responseCoding, train_y, cv_
```


Log loss : 1.0732328472930994


Number of mis-classified points : 0.3815789473684211

----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----





----- Recall matrix (Row sum=1) -----



4.3. Logistic Regression

4.3.1. With Class balancing

4.3.1.1. Hyper paramter tuning

In [116]:

```
# read more about SGDClassifier() at http://scikit-learn.org/stable/mo
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.1
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, l
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stoc
# predict(X) Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.o
# -----
# default paramters
```

```

# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, metho
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-5, 1)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2')
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf
    # to avoid rounding error while multiplying probabilities we use lo
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha],
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross val
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log


for alpha = 1e-05
Log Loss : 1.0820257983988437
for alpha = 0.0001
Log Loss : 1.0264629789850206
for alpha = 0.001
Log Loss : 0.9618451656833001

```

```

for alpha = 0.01
Log Loss : 1.0136912392422004
for alpha = 0.1
Log Loss : 1.1425637431283917
for alpha = 1
Log Loss : 1.3763226423522668

```



```

For values of best alpha = 0.001 The train log loss is: 0.525758897
0311815
For values of best alpha = 0.001 The cross validation log loss is:
0.9618451656833001
For values of best alpha = 0.001 The test log loss is: 0.9183512118
693308

```

4.3.1.2. Testing the model with best hyper paramters

```

In [117]: # read more about SGDClassifier() at http://scikit-learn.org/stable/mo
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.1,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, l
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stoc
# predict(X) Predict class labels for samples in X.

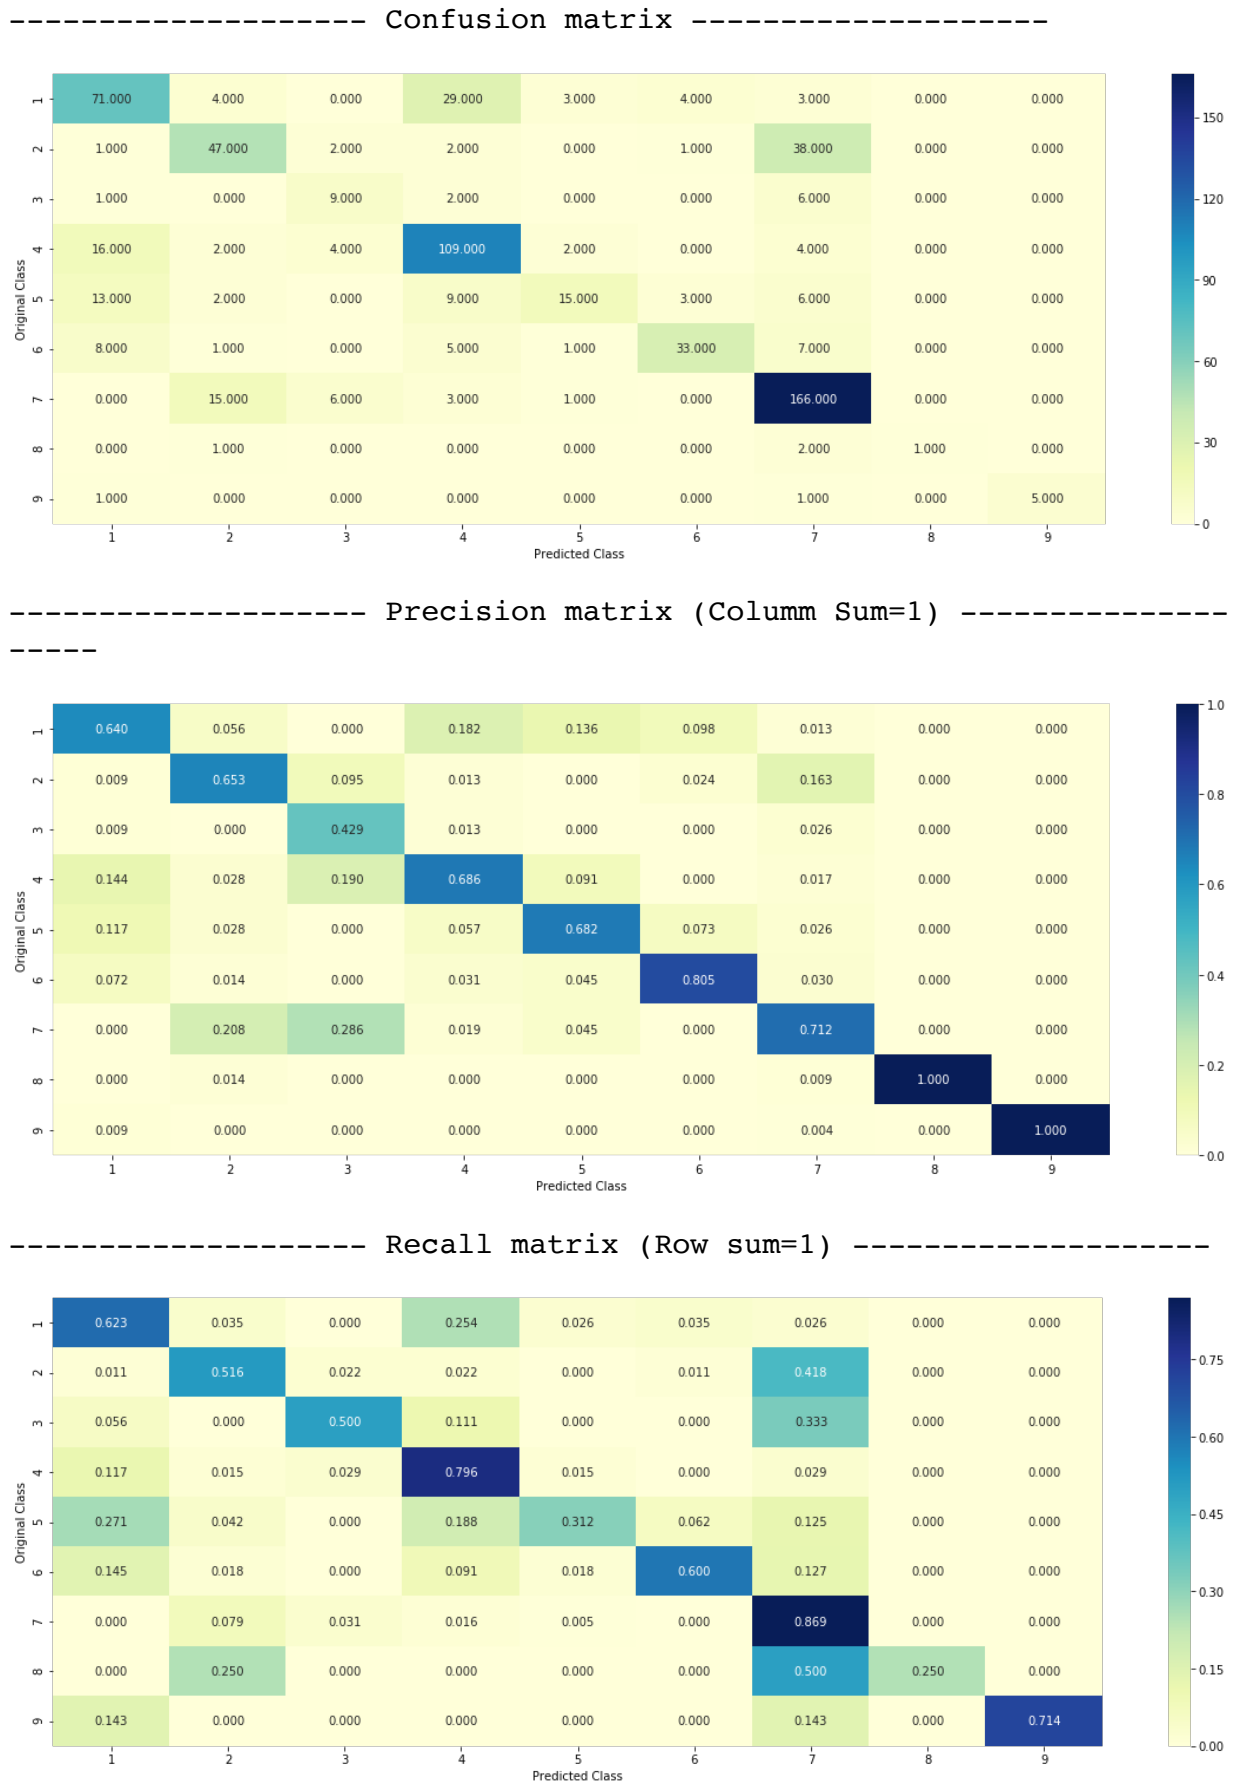
#-----
# video link: https://www.appliedaiaicourse.com/course/applied-ai-course
#-----
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha],
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y, test_y

```

```

Log loss : 0.9183512118693308
Number of mis-classified points : 0.3142857142857143

```



4.3.1.3. Feature Importance

```
In [118]: def get_imp_feature_names(text, indices, removed_ind = []):
word_present = 0
tabulte_list = []
incresingorder_ind = 0
for i in indices:
    if i < train_gene_feature_onehotCoding.shape[1]:
        tabulte_list.append([incresingorder_ind, "Gene", "Yes"])
    elif i < 18:
        tabulte_list.append([incresingorder_ind, "Variation", "Yes"])
    if ((i > 17) & (i not in removed_ind)) :
        word = train_text_features[i]
        yes_no = True if word in text.split() else False
        if yes_no:
            word_present += 1
            tabulte_list.append([incresingorder_ind, train_text_features[i], "Yes"])
        incresingorder_ind += 1
print(word_present, "most important features are present in our query")
print("-"*50)
print("The features that are most important of the ", predicted_cls)
print(tabulate(tabulte_list, headers=["Index", "Feature name", "Present"])
```

4.3.1.3.1. Correctly Classified point

```
In [119]: # from tabulate import tabulate
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], random_state=1)
clf.fit(train_x_onehotCoding, train_y)
test_point_index = 1
no_feature = 50
predicted_cls = sig_clf.predict(test_x_onehotCoding.iloc[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotCoding.iloc[test_point_index])[0], 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
print("-"*50)
```

```
Predicted Class : 1
Predicted Class Probabilities: [[0.4303 0.1552 0.0084 0.1265 0.011
0.0039 0.0196 0.009 0.236 ]]
Actual Class : 1
-----
```

4.3.1.3.2. correctly Classified point

```
In [120]: # from tabulate import tabulate
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha],
clf.fit(train_x_onehotCoding,train_y)
test_point_index = 42
no_feature = 50
predicted_cls = sig_clf.predict(test_x_onehotCoding.iloc[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotCoding.iloc[test_point_index]), 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
print("-"*50)
```

```
Predicted Class : 1
Predicted Class Probabilities: [[0.7581 0.1618 0.0083 0.0175 0.0191
0.0159 0.0107 0.0061 0.0025]]
Actual Class : 1
-----
```

4.3.2. Without Class balancing

4.3.2.1. Hyper paramter tuning

```
In [121]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='isotonic')
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
"
```

```

#-----

alpha = [10 ** x for x in range(-3, 1)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=0)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_))
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=0)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)


predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is: ", log_loss(train_y, predict_y))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is: ", log_loss(cv_y, predict_y))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is: ", log_loss(test_y, predict_y))


```

```

for alpha = 0.001
Log Loss : 0.9549254257049363
for alpha = 0.01
Log Loss : 0.9933906777381515
for alpha = 0.1
Log Loss : 1.1443484005793763
for alpha = 1
Log Loss : 1.404725094137652

```





For values of best alpha = 0.001 The train log loss is: 0.5168167510834673

For values of best alpha = 0.001 The cross validation log loss is: 0.9549254257049363

For values of best alpha = 0.001 The test log loss is: 0.9125467886588063

4.3.2.2. Testing model with best hyper parameters

```
In [122]: # read more about SGDClassifier() at http://scikit-learn.org/stable/mo
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.1
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, l
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stoc
# predict(X) Predict class labels for samples in X.


#-----
# video link:
#-----



clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y, test_
```

Log loss : 0.9125467886588063

Number of mis-classified points : 0.31278195488721805

----- Confusion matrix -----





4.3.2.3. Feature Importance, Correctly Classified point

```
In [123]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
clf.fit(train_x_onehotCoding,train_y)
test_point_index = 42
no_feature = 50
predicted_cls = sig_clf.predict(test_x_onehotCoding.iloc[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotCoding.iloc[test_point_index]), 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
print("-"*50)
```

```
Predicted Class : 1
Predicted Class Probabilities: [[0.7494 0.175  0.0074 0.0172 0.02
0.0158 0.0089 0.0038 0.0024]]
Actual Class : 1
-----
```

4.3.2.4. Feature Importance, Inorrectly Classified point

```
In [124]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
clf.fit(train_x_onehotCoding,train_y)
test_point_index = 21
no_feature = 50
predicted_cls = sig_clf.predict(test_x_onehotCoding.iloc[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotCoding.iloc[test_point_index]), 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
print("-"*50)
```

```
Predicted Class : 7
Predicted Class Probabilities: [[0.001  0.1837 0.0062 0.0092 0.0257
0.0216 0.7443 0.0041 0.0043]]
Actual Class : 2
-----
```

4.4. Linear Support Vector Machines

4.4.1. Hyper paramter tuning

```
In [125]: # read more about support vector machines with linear kernals here http://www.cs.cmu.edu/~tom/10-708/lect07.pdf
# -----
# default parameters
# SVC(C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True, cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_scaling=1)
# Some of methods of SVM()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data X and target values y. If sample_weight is provided, then the loss function is weighted by the sample_weight.
```

```

# fit(X, y, [sample_weight])    Fit the SVM model according to the given
# predict(X)    Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/
# -----
# default parameters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid')
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])    Fit the calibrated model
# get_params([deep])    Get parameters for this estimator.
# predict(X)    Predict the target of new samples.
# predict_proba(X)    Posterior probabilities of classification
# -----
# video link:
# -----

alpha = [10 ** x for x in range(-5, 3)]
cv_log_error_array = []
for i in alpha:
    print("for C =", i)
    # clf = SVC(C=i, kernel='linear', probability=True, class_weight='balanced')
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2')
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_))
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
# clf = SVC(C=i, kernel='linear', probability=True, class_weight='balanced')
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2')
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)

```

```


print('For values of best alpha = ', alpha[best_alpha], "The train log
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross val
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log

```

```

for C = 1e-05
Log Loss : 1.0709184053649463
for C = 0.0001
Log Loss : 1.090737090452236
for C = 0.001
Log Loss : 1.0516495767011695
for C = 0.01
Log Loss : 1.0529856698264046
for C = 0.1
Log Loss : 1.0835262662532756
for C = 1
Log Loss : 1.077480200178823
for C = 10
Log Loss : 1.5065398393210485
for C = 100
Log Loss : 1.8034325767071846

```



```

For values of best alpha = 0.001 The train log loss is: 0.572901793
1050538
For values of best alpha = 0.001 The cross validation log loss is:
1.0516495767011695
For values of best alpha = 0.001 The test log loss is: 0.9975590275
165219

```

4.4.2. Testing model with best hyper parameters

```

In [126]: # read more about support vector machines with linear kernal here http
# -----
# default parameters

```

```
# SVC(C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinkin
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decis


# Some of methods of SVM()
# fit(X, y, [sample_weight])    Fit the SVM model according to the giv
# predict(X)    Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
# -----

# clf = SVC(C=alpha[best_alpha],kernel='linear',probability=True, clas
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y,cv_x_on
```


Log loss : 1.0516495767011695

Number of mis-classified points : 0.34022556390977443


----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----





4.5 Random Forest Classifier

4.5.1. Hyper paramter tuning (With One hot Encoding)

```
In [128]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])    Fit the SVM model according to the giv
# predict(X)    Perform classification on samples in X.
# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature)

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.o
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, metho
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])    Fit the calibrated model
# get_params([deep])    Get parameters for this estimator.
# predict(X)    Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
```

```

#-----

alpha = [100,200,500,1000,2000]
max_depth = [5, 10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators =", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini',
        clf.fit(train_x_onehotCoding, train_y)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_x_onehotCoding, train_y)
        sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels)
        print("Log Loss :",log_loss(cv_y, sig_clf_probs))

'''fig, ax = plt.subplots()
features = np.dot(np.array(alpha)[: ,None],np.array(max_depth)[None]).r
ax.plot(features, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[int(i/2)],max_depth[int(i%2)],str(txt)), (featu
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
'''

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], cr
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The

for n_estimators = 100 and max depth = 5
Log Loss : 1.221251856458788
for n_estimators = 100 and max depth = 10
Log Loss : 1.1523133259436458
for n_estimators = 200 and max depth = 5
Log Loss : 1.215983059461481
for n_estimators = 200 and max depth = 10
Log Loss : 1.1445768085235268
for n_estimators = 500 and max depth = 5
Log Loss : 1.2074140678407117
for n_estimators = 500 and max depth = 10
Log Loss : 1.1396000410342781
for n_estimators = 1000 and max depth = 5

```



```

Log Loss : 1.1995699341288346
for n_estimators = 1000 and max depth = 10
Log Loss : 1.1379069861482
for n_estimators = 2000 and max depth = 5
Log Loss : 1.1956724884838927
for n_estimators = 2000 and max depth = 10
Log Loss : 1.1327758191013155
For values of best estimator = 2000 The train log loss is: 0.629269
5980731973
For values of best estimator = 2000 The cross validation log loss i
s: 1.1327758191013153
For values of best estimator = 2000 The test log loss is: 1.0772846
773426699

```

4.5.2. Testing model with best hyper parameters (One Hot Encoding)

```

In [129]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight]) Fit the SVM model according to the giv
# predict(X) Perform classification on samples in X.
# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature)

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
# -----


clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], cr
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y,cv_x_on


```

```


Log loss : 1.1327758191013153
Number of mis-classified points : 0.36278195488721804
----- Confusion matrix -----

```






----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.5.3. Hyper paramter tuning (With Response Coding)

```
In [130]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])    Fit the SVM model according to the giv
# predict(X)    Perform classification on samples in X.
```

```

# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature)

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.o
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, metho
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10,50,100,200,500,1000]
max_depth = [2,3,5,10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators =", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini',
        clf.fit(train_x_responseCoding, train_y)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_x_responseCoding, train_y)
        sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels
        print("Log Loss :",log_loss(cv_y, sig_clf_probs))
    ...

fig, ax = plt.subplots()
features = np.dot(np.array(alpha)[: ,None],np.array(max_depth)[None]).r
ax.plot(features, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[int(i/4)],max_depth[int(i%4)],str(txt)), (featu
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
'''

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/4)], cr
clf.fit(train_x_responseCoding, train_y)

```

```

sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The tr
predict_y = sig_clf.predict_proba(cv_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The cr
predict_y = sig_clf.predict_proba(test_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The te

```

```

for n_estimators = 10 and max depth = 2
Log Loss : 2.985604279662273
for n_estimators = 10 and max depth = 3
Log Loss : 2.5805638925724144
for n_estimators = 10 and max depth = 5
Log Loss : 2.50874108489422
for n_estimators = 10 and max depth = 10
Log Loss : 2.500396140877808
for n_estimators = 50 and max depth = 2
Log Loss : 2.288155547931257
for n_estimators = 50 and max depth = 3
Log Loss : 2.325426838395101
for n_estimators = 50 and max depth = 5
Log Loss : 2.1372277705559157
for n_estimators = 50 and max depth = 10
Log Loss : 2.3333314077055536
for n_estimators = 100 and max depth = 2
Log Loss : 2.372609968962405
for n_estimators = 100 and max depth = 3
Log Loss : 2.557601374610918
for n_estimators = 100 and max depth = 5
Log Loss : 2.096004993001543
for n_estimators = 100 and max depth = 10
Log Loss : 2.346585959841944
for n_estimators = 200 and max depth = 2
Log Loss : 2.403972042361107
for n_estimators = 200 and max depth = 3
Log Loss : 2.5859570133876453
for n_estimators = 200 and max depth = 5
Log Loss : 2.251926017679841
for n_estimators = 200 and max depth = 10
Log Loss : 2.4353192227554064
for n_estimators = 500 and max depth = 2
Log Loss : 2.4422391006300255
for n_estimators = 500 and max depth = 3
Log Loss : 2.488401064217581
for n_estimators = 500 and max depth = 5
Log Loss : 2.2711479592792223
for n_estimators = 500 and max depth = 10
Log Loss : 2.4646423499305707
for n_estimators = 1000 and max depth = 2
Log Loss : 2.4038788188197193
for n_estimators = 1000 and max depth = 3
Log Loss : 2.34861746224742

```

```

for n_estimators = 1000 and max depth = 5
Log Loss : 2.216408979041495
for n_estimators = 1000 and max depth = 10
Log Loss : 2.4344455062889496
For values of best alpha = 100 The train log loss is: 0.02427030396
5366887
For values of best alpha = 100 The cross validation log loss is: 2.
096004993001543
For values of best alpha = 100 The test log loss is: 2.067159216343
359

```

4.5.4. Testing model with best hyper parameters (Response Coding)

```

In [131]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight]) Fit the SVM model according to the giv
# predict(X) Perform classification on samples in X.
# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature)

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
# -----

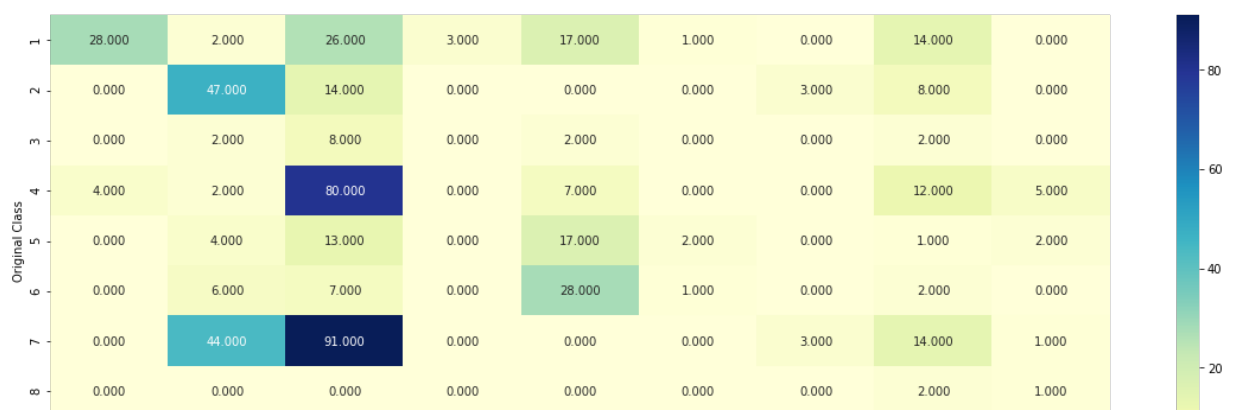
clf = RandomForestClassifier(max_depth=max_depth[int(best_alpha%4)], n
predict_and_plot_confusion_matrix(train_x_responseCoding, train_y,cv_x


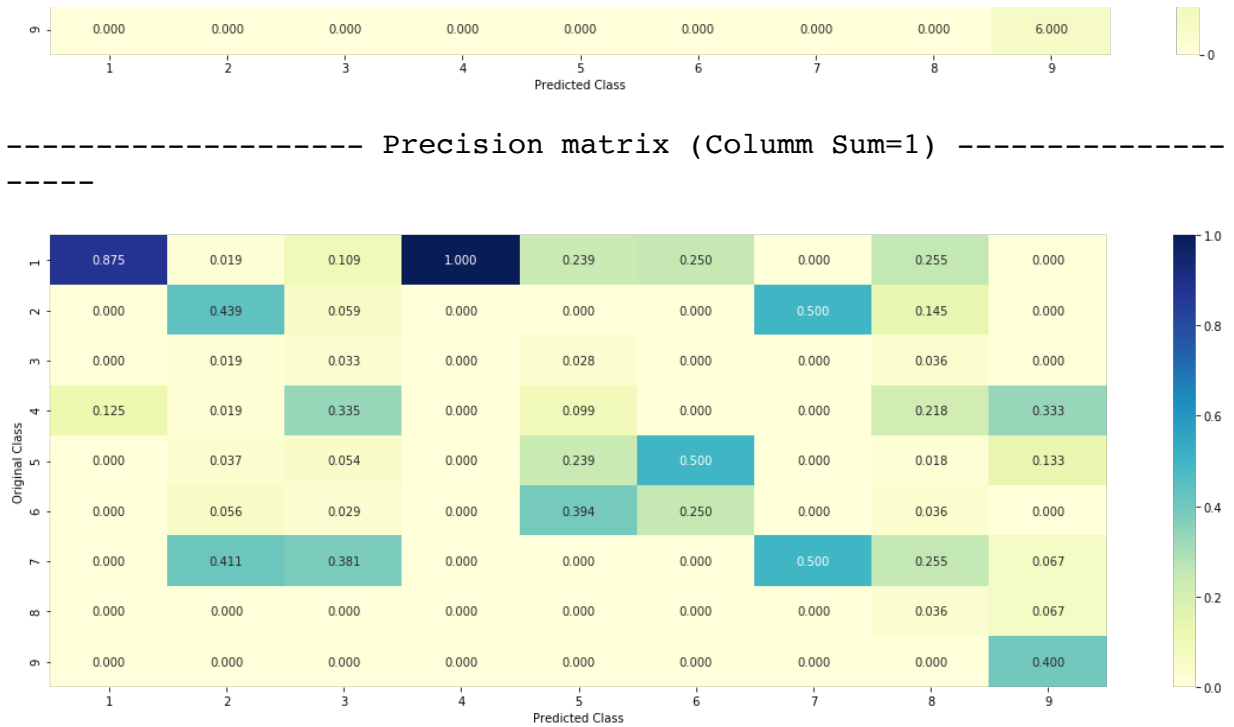
```

Log loss : 2.096004993001543

Number of mis-classified points : 0.7894736842105263

----- Confusion matrix -----





4.7 Stack the models

4.7.1 testing with hyper parameter tuning

```
In [133]: # read more about SGDClassifier() at http://scikit-learn.org/stable/mo
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.1
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, l
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stoc
```

```

# predict(X)      Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
#-----

# read more about support vector machines with linear kernels here http://www.appliedaicourse.com/course/applied-ai-course
# -----
# default parameters
# SVC(C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True, cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='raw')

# Some of methods of SVM()
# fit(X, y, [sample_weight])    Fit the SVM model according to the given training data.
# predict(X)      Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
# -----

# read more about support vector machines with linear kernels here http://www.appliedaicourse.com/course/applied-ai-course
# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])    Fit the SVM model according to the given training data.
# predict(X)      Perform classification on samples in X.
# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature)

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course
# -----

clf1 = SGDClassifier(alpha=0.001, penalty='l2', loss='log', class_weight=None)
clf1.fit(train_x_onehotCoding, train_y)
sig_clf1 = CalibratedClassifierCV(clf1, method="sigmoid")

clf2 = SGDClassifier(alpha=1, penalty='l2', loss='hinge', class_weight=None)
clf2.fit(train_x_onehotCoding, train_y)
sig_clf2 = CalibratedClassifierCV(clf2, method="sigmoid")

clf3 = MultinomialNB(alpha=0.001)
clf3.fit(train_x_onehotCoding, train_y)

```

```

sig_clf3 = CalibratedClassifierCV(clf3, method="sigmoid")

sig_clf1.fit(train_x_onehotCoding, train_y)
print("Logistic Regression : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf1.predict_proba(cv_x_onehotCoding))))
sig_clf2.fit(train_x_onehotCoding, train_y)
print("Support vector machines : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf2.predict_proba(cv_x_onehotCoding))))
sig_clf3.fit(train_x_onehotCoding, train_y)
print("Naive Bayes : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf3.predict_proba(cv_x_onehotCoding))))
print("-"*50)
alpha = [0.0001,0.001,0.01,0.1,1,10]
best_alpha = 999
for i in alpha:
    lr = LogisticRegression(C=i)
    sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3], method="sigmoid")
    sclf.fit(train_x_onehotCoding, train_y)
    print("Stacking Classifier : for the value of alpha: %f Log Loss: %f" % (i, log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))))
    log_error = log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))
    if best_alpha > log_error:
        best_alpha = log_error

```

```

Logistic Regression : Log Loss: 0.95
Support vector machines : Log Loss: 1.09
Naive Bayes : Log Loss: 1.22
-----
Stacking Classifier : for the value of alpha: 0.000100 Log Loss: 2.175
Stacking Classifier : for the value of alpha: 0.001000 Log Loss: 2.010
Stacking Classifier : for the value of alpha: 0.010000 Log Loss: 1.431
Stacking Classifier : for the value of alpha: 0.100000 Log Loss: 1.048
Stacking Classifier : for the value of alpha: 1.000000 Log Loss: 1.081
Stacking Classifier : for the value of alpha: 10.000000 Log Loss: 1.308

```

4.7.2 testing the model with the best hyper parameters

```

In [134]: lr = LogisticRegression(C=0.1)
sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3], method="sigmoid")
sclf.fit(train_x_onehotCoding, train_y)

log_error = log_loss(train_y, sclf.predict_proba(train_x_onehotCoding))
print("Log loss (train) on the stacking classifier :",log_error)

log_error = log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))
print("Log loss (CV) on the stacking classifier :",log_error)

log_error = log_loss(test_y, sclf.predict_proba(test_x_onehotCoding))
print("Log loss (test) on the stacking classifier :",log_error)

```



```
print('Number of missclassified point :', np.count_nonzero(sclf.predict(
plot_confusion_matrix(test_y=test_y, predict_y=sclf.predict(test_x_one)
```


Log loss (train) on the stacking classifier : 0.6013716058609647

Log loss (CV) on the stacking classifier : 1.0475754062476912


Log loss (test) on the stacking classifier : 1.0445581939340607

Number of missclassified point : 0.35037593984962406


----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----





4.7.3 Maximum Voting classifier

```
In [135]: #Refer:http://scikit-learn.org/stable/modules/generated/sklearn.ensembl
from sklearn.ensemble import VotingClassifier
vclf = VotingClassifier(estimators=[('lr', sig_clf1), ('svc', sig_clf2)
vclf.fit(train_x_onehotCoding, train_y)
print("Log loss (train) on the VotingClassifier :", log_loss(train_y, vclf.predict(train_x_onehotCoding)))
print("Log loss (CV) on the VotingClassifier :", log_loss(cv_y, vclf.predict(cv_x_onehotCoding)))
print("Log loss (test) on the VotingClassifier :", log_loss(test_y, vclf.predict(test_x_onehotCoding)))
print("Number of missclassified point :", np.count_nonzero(test_y - vclf.predict(test_x_onehotCoding)))
plot_confusion_matrix(test_y, vclf.predict(test_x_onehotCoding))
```


Log loss (train) on the VotingClassifier : 0.6997619720475359

Log loss (CV) on the VotingClassifier : 1.0295222706087999

Log loss (test) on the VotingClassifier : 1.0155378435687559

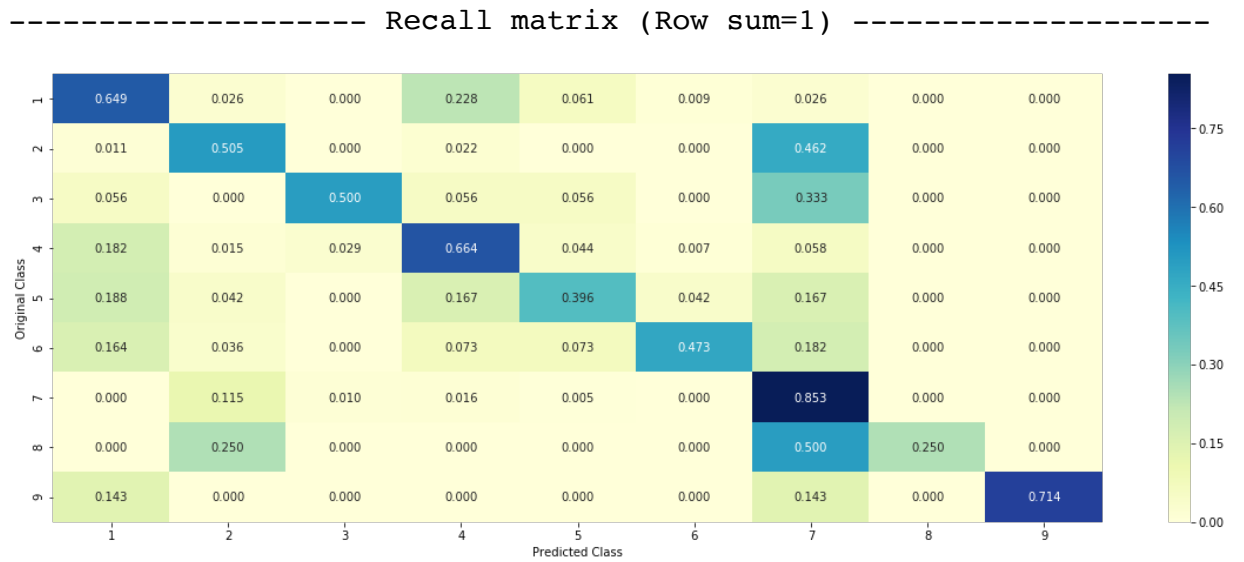
Number of missclassified point : 0.3473684210526316

----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----





```
In [136]: from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Model", "Train_logloss", "Cv_logloss", "Test_logloss",
                 "Misclassified_error"]

x.add_row(["Naive Bayes", 0.79453, 1.21285, 1.23219, 35.52])
x.add_row(["KNN", 0.75710, 1.07323, 1.09839, 38.15])
x.add_row(["LR (Balanced data)", 0.52575, 0.96184, 0.91835, 31.42])
x.add_row(["LR (Without Balanced data)", 0.51681, 0.95492, 0.91254, 31.27])
x.add_row(["Linear SVM", 0.57290, 1.05164, 0.99755, 34.02])
x.add_row(["Random Forest (one hot encoding)", 0.62926, 1.13277, 1.07728, 36.27])
x.add_row(["Random Forest (Response Coding)", 0.02427, 2.09600, 2.06715, 78.94])
x.add_row(["Stacking model (LR, SVM, NB)", 0.60137, 1.04757, 1.04455, 35.03])
x.add_row(["Maximum Voting Classifier", 0.69976, 1.02952, 1.01553, 34.73])

print(x)
```

Model	Train_logloss	Cv_logloss	Test_logloss	Misclassified_error
Naive Bayes	0.79453	1.21285	1.23219	35.52
KNN	0.7571	1.07323	1.09839	38.15
LR (Balanced data)	0.52575	0.96184	0.91835	31.42
LR (Without Balanced data)	0.51681	0.95492	0.91254	31.27
Linear SVM	0.5729	1.05164	0.99755	34.02
Random Forest (one hot encoding)	0.62926	1.13277	1.07728	36.27
Random Forest (Response Coding)	0.02427	2.096	2.06715	78.94
Stacking model (LR, SVM, NB)	0.60137	1.04757	1.04455	35.03
Maximum Voting Classifier	0.69976	1.02952	1.01553	34.73

Steps followed

- 1) For reducing the loss in test data, Featurization for the text data used was TF-IDF Featurization. The least loss score for train, test, CV was using Logistic Regression algorithm and class being balanced.
- 2) Instead of using all the features, the top 1000 features in TF-IDF features were selected. The least loss score for train, test, CV was using Logistic Regression algorithm and class being balanced. KNN algorithm also performed well but the model is not interpretable and interpretability is required in this given problem.
- 3) As Logistic regression (both balanced and unbalanced class) was performing better than other algorithms. Uni-gram and Bi-gram in Bag of Words featurization was used. The least loss score for train, test, CV was using Logistic Regression algorithm and class being balanced. Linear SVM and KNN algorithm also performed well when compared to other featurization. This also could not reduce the log loss less than 1.
- 4) To get log loss less than 1, A few featurization techniques were implemented.
 - 1) Number of words in a given column.
 - 2) Number of characters in a given column.
 - 3) Combining of Gene feature and variation in a single column.
 - 4) Number of words greater than 5000 in a given column. If it is greater than the value assigned is 1 else the value assigned is 0.
 - 5) Number of characters greater than 50000 in a given column. If it is greater than the value is 1 else the values assigned is 0.

The Log loss was less than 1 using Logistic regression with both balanced and imbalanced class.

5. Observation

To get log loss of less than 1.0, a few feature engineering was implemented :

- 1) Number of words in a given column
- 2) Number of characters in a given column
- 3) Combing gene and variation in a single column
- 4) Number of words greater than 5000 in a column, if true then 1 else 0
- 5) Number of chracters greater than 50000 in a column, if true then 1 else 0

The log loss is less than 1 for Logistic Regression with both balanced data and imbalanced data